

The Sigma Manual
A Guide for Users and Developers
of the SUMO Toolset

Adam Pease

March 28, 2018

Contents

1	Preface	7
1.1	Introduction	7
2	Code	11
2.1	Hello World	11
2.2	Coding Standards	13
2.3	Operational Sequence	14
2.4	Configuration	16
2.5	Conversion of SUO-KIF to First-Order Logic	20
2.6	Normalization Algorithm	26
2.7	Conversion to TPTP	28
3	Using Sigma	31
3.1	JSP Interface	31
3.2	Language Generation	34
3.3	Browsing and Display	37
3.4	Analysis and Debugging	39
3.5	Inference	40
3.6	Proof Presentation	43
3.7	Mapping, Merging and Translation	44
3.8	Working with Sigma and SUMO	45
3.9	Browsing and Editing	46
3.10	Sigma Tutorial	47
4	Conclusions	61
4.1	Acknowledgements	61

List of Figures

1.1	Major Sigma functions	8
2.1	CodeFormat1	13
2.2	CodeFormat2	14
2.3	Example axioms showing predicate variables	20
2.4	Example axioms with holds prefixing	21
2.5	"holds" insertion	21
2.6	Holds prefix with arity	22
2.7	Example axiom with predicate instantiation	22
2.8	Axiom with row variables	23
2.9	Row variable expansion	23
2.10	Example axiom	25
2.11	Sortal prefixing	25
2.12	Naive sortal prefixing	26
3.1	Sigma browser screen	35
3.2	An axiom for Ambulating	37
3.3	Simplified browser view	38
3.4	Proof presentation in Sigma, from [Trac et al., 2008]	41
3.5	New query, negating the first result in order to find a new answer.	42
3.6	Example contradictory knowledge base	42
3.7	Example redundant knowledge base	44
3.8	Sigma login screen	47
3.9	The Knowledge Bases screen	48
3.10	Manifest	49
3.11	Manifest2	50
3.12	Initial term browser screen	51
3.13	Term browser page for the SUMO term Walking	52
3.14	Term neighbors	53
3.15	English word listing	54
3.16	Textual Graph page	55
3.17	Graph page	56
3.18	Sigma proof	57
3.19	Preferences page	59

Chapter 1

Preface

1.1 Introduction

The Suggested Upper Merged Ontology (SUMO) began as a simple project to create a high-level taxonomic structure for computer assisted training applications, but it quickly became more ambitious, attempting to cover any topics for any application that addressed entities in our common-sense world, and that were amenable to symbolic processing. It grew to cover a wide range of representational constructs, challenging the practical use of research in knowledge representation and reasoning.

The development of SUMO also came to encompass tools used in its development and application. The Sigma system [Pease, 2003] began as a simple ontology browsing aid but expanded to incorporate reasoning tools created by others, first the Vampire theorem prover [Riazanov and Voronkov, 2002] and then an entire suite of reasoners called TPTP [Trac et al., 2008], especially the E theorem prover [Schulz, 2002]. That work in turn required efforts on making reasoning practical and efficient on such a rich knowledge base [Pease et al., 2010].

A second line of work was in relating SUMO to linguistic data, first in developing the links from the WordNet [Fellbaum, 1998] lexical database [Niles and Pease, 2003] and then developing a natural language understanding system, called the Controlled English to Logic Translation system (CELT) that used those links, among much other information, to translate language to logic [Pease and Li, 2010]. That work has been superseded by an approach called Semantic Rewriting, which is embodied in the SigmaNLP system, whose manual is forthcoming.

The Sigma Knowledge Engineering Environment is a development environment for logical theories, specifically those that extend the Suggested Upper Merged Ontology (SUMO). Sigma is analogous to an Integrated Development Environment (IDE) for programming languages like C++ or Java. Modern and expressive languages for the development of formal theories, such as SUO-KIF [Pease, 2009] and TPTP [Sutcliffe, 2009] have a similar degree of expressiveness,

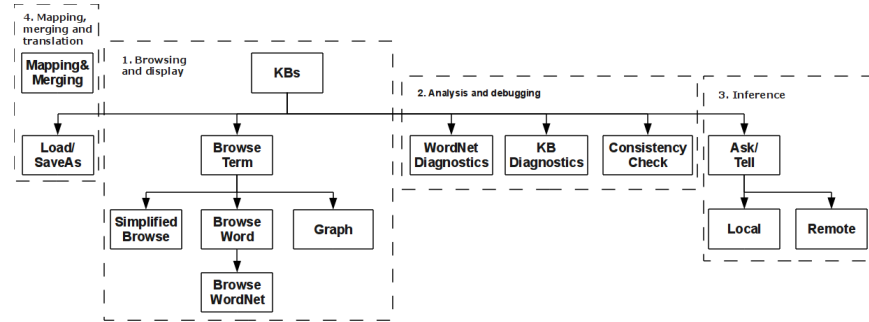


Figure 1.1: Major Sigma functions

in a broad sense, to a modern programming language. Development of SUMO is done in a text editor and Sigma provides support for browsing, debugging and applying SUMO to applications. Like all general-purpose programming languages, it is too expressive to allow development in any sort of purely visual mode. This makes it distinct from taxonomy languages, or description logics, in which the languages are defined primarily by their graph structure.

We have utilized the Git for collaborative ontology development rather than implementing some additional tool for collaboration. Developers are typically given authority over one or more ontologies, required to check in progress at least weekly so that other developers can sync up with their changes. This has also resulted in a detailed public record of the development and evolution of SUMO [Pease and Benzmler, 2010].

Ontology is a broad area, and there are many approaches to ontology. The purpose of this manual is not to teach ontology, or SUMO, so a prerequisite for starting to work with Sigma is to read the book "Ontology: A Practical Guide", and complete all the exercises in the book. Otherwise, many aspects of Sigma will be mysterious.

This manual is a prerequisite for understanding the functionality of the SigmaNLP system, although that manual is yet to be written. My hope is that this manual will be a "living" document, subject to frequent updating as Sigma is improved and extended. At the moment, Sigma has a number of components that are obsolete, or in the process of being completed or redone.

While Sigma [Pease, 2013] [Pease, 2003] was created to support SUMO, and that has been its primary use during some eight years of development, that is by no means the only theory that it can handle. Sigma works on knowledge bases that can be composed from various files selected by the user. Those files can be coded in a small number of different formal languages, including TPTP and OWL, as well as SUO-KIF. The Sigma user can easily work with very small theories or very large ones by composing only the theories that are needed for the work at hand. A typical use of Sigma would involve loading just the upper level of SUMO and whatever domain theory is needed for the user's chosen application area.

Tools within Sigma (Figure 1.1) can be broadly segmented into several groups, (1) browsing and display, (2) analysis and debugging, (3) inference, and (4) mapping, merging and translation.

In this manual we provide an index (starting on page 65) to all the classes mentioned, so that it can be used as a reference going beyond what one can find in the JavaDoc for Sigma¹

¹which is online at <http://www.ontologyportal.org/sigma-keel-doc/>

Chapter 2

Code

2.1 Hello World

The best way to learn how to use a code library is to start doing something with it. So, we'll start with a "Hello World" example for using Sigma. We'll keep it simple and just use the command line. We'll assume use on Linux as well.

First, follow the installation instructions for Sigma at <https://github.com/ontologyportal/sigma>. Most problems in installation come from people rushing through them too quickly, without reading the documentation in full, or skipping a step by mistake, or trying to alter the steps without fully understanding the implications of a change. So, try installing exactly as instructed first, then maybe uninstall and try again, altering things to your taste, if that's essential. Make sure you can complete the instructions and run Sigma before trying the "Hello World".

We'll try just initializing Sigma and asking what is the parent class of the SUMO term `PrimaryColor`.

```
import com.articulate.sigma.*;

public class MySigma {

    public static void main(String[] args) {

        KBmanager.getMgr().initializeOnce();
        KB kb = KBmanager.getMgr().getKB("SUMO");
        System.out.println(kb.immediateParents("PrimaryColor"));
    }
}
```

The first line initializes Sigma, causing it to read the `config.xml` file in your `$SIGMA_HOME/KBs` directory. That causes Sigma to set all the parameters specified in that file, and then read all the knowledge bases it lists. If the knowledge base files have been read previously, and not changed, then Sigma can quickly load all the data from a serialized file. If it has to load all the

ontology files from their original, textual source, the process takes much longer. Sigma will also read the WordNet and Open Multilingual Wordnet lexicons and several supporting files.

The next line gets the combined knowledge base named in the config.xml file.

The last line prints out all the parents of `PrimaryColor`.

Compile with

```
javac -classpath $SIGMA_SRC/build/*:$SIGMA_SRC/lib/* MySigma.java
```

and run with

```
java -classpath $SIGMA_SRC/build/*:$SIGMA_SRC/lib/*:. MySigma
```

```
Info in KBmanager.initializeOnce()
Info in KBmanager.initializeOnce(): initializing with
/home/apease/.sigmakee/KBs
KBmanager.readConfiguration()
KBmanager.serializedExists(): true
KBmanager.serializedOld(config):
KBmanager.serializedOld(config): save date: Wed Mar 14
10:39:20 PDT 2018
kbsFileNamesFromXML(): Completed loading KB names
KBmanager.serializedOld(config): returning false (not old)
KBmanager.loadSerialized(): KBmanager has been deserialized
WordNet.initOnce(): using baseDir =
/home/apease/.sigmakee/KBs/WordNetMappings
WordNet.loadSerialized(): WN has been deserialized
ENTER DB.readSpreadsheet(/home/apease/.sigmakee/KBs/
WordNetMappings/sentiment.csv, null)
ENTER DB.readSpreadsheet(java.io.FileReader@28ba21f3, null)
EXIT DB.readSpreadsheet(java.io.FileReader@28ba21f3, null)
rows == [list of 8222 rows]
0.313 seconds elapsed time
EXIT DB.readSpreadsheet(/home/apease/.sigmakee/KBs/
WordNetMappings/sentiment.csv, null)
NLGUutils.init(): initializing with /home/apease/.sigmakee/KBs
NLGUutils.readKeywordMap():
NLGUutils.serializedExists(): true
NLGUutils.loadSerialized()
NLGUutils.loadSerialized(): NLGUutils has been deserialized
INFO in OMWordnet.readOMWfiles(): reading files:
OMWordnet.loadSerialized(): OMW has been deserialized
Info in KBmanager.initializeOnce(): initialized is true
[ColorAttribute]
```

If all goes well, you should see output very similar to this, ending with the resulting list of parent terms of `PrimaryColor`.

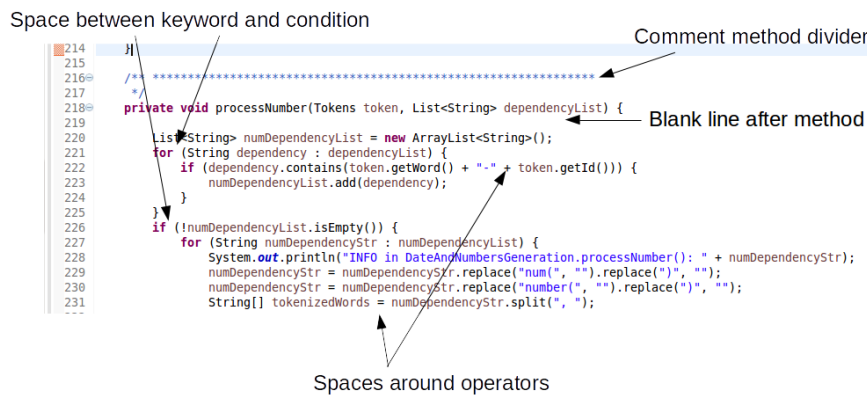


Figure 2.1: CodeFormat1

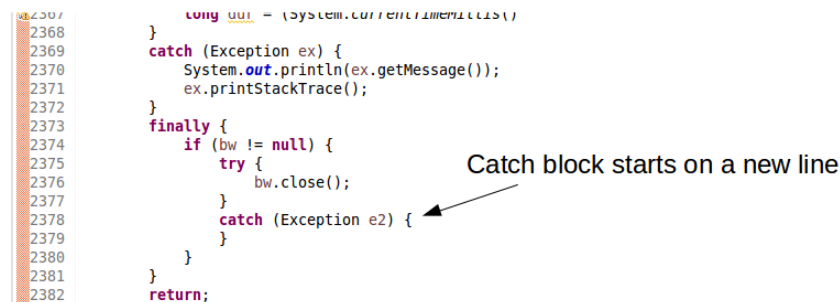
2.2 Coding Standards

The watchword for Sigma development is "simplicity" (although this may not be apparent to a new user!). After 18 years of effort, at the time of this writing, we aim for long term development. This means cautious use of external tools or libraries or even Java language features that might result in dependencies that could change frequently. It also means doing development to fulfill a particular needed function, without creating things that just might be needed for the future. In this way, we subscribe to the "You Ain't Gonna Need It" mantra of eXtreme Programming.

Many Sigma classes have a command line interface. This is useful for testing and also helps developers by providing minimal, lightweight interfaces to various utilities. This has been used much more in SigmaNLP but we're attempting to do more of this in SigmaKEE. We attempt to follow a Unix style of command invocation. Look at **KBmanager** for a model.

Another imprecisely-realized goal is test-driven development. Ideally, every significant method will have a junit test. We employ a convention of dividing tests into three groups. Unit tests are relatively quick and should always pass. Integration tests may be considerably slower, and involve operations on large knowledge bases, but should always pass. Corpus tests may be slow and may not pass, but indicate a degree of progress on research goals that may not be solved yet. Tests are kept in a `text` directory in both SigmaKEE and SigmaNLP.

Sigma has a standard code formatting throughout, which is especially important given the long timeframe of development and the fact that many people have contributed to this open source product. While any given feature may not be to any programmer's liking, consistency overall is very important to the code's readability. Figures 2.1 and 2.2 illustrate the contentions used.



```

2367     long uid = (System.currentTimeMillis()
2368     }
2369     catch (Exception ex) {
2370         System.out.println(ex.getMessage());
2371         ex.printStackTrace();
2372     }
2373     finally {
2374         if (bw != null) {
2375             try {
2376                 bw.close();
2377             }
2378             catch (Exception e2) {
2379             }
2380         }
2381     }
2382     return;

```

Figure 2.2: CodeFormat2

2.3 Operational Sequence

KBmanager is a class that handles dispatching the various tasks that initialize Sigma. It fetches the config.xml file that controls the parameters set for Sigma, and the knowledge base files that are loaded. It also loads all of the WordNet lexicon, the semantic links in WordNet, and the links from WordNet to SUMO. It loads the Open Multilingual Wordnet that has languages other than English. **KBmanager** is **Serializable**, since there are a number of files to load, and the operations on the knowledge base files can take a while, so being able to save the files and their transformations and indexes as binary can greatly speed loading time. **KBmanager** also has a simple hook for a Python interface in the method `pythonServer()`, which can be started by invoking **KBmanager** from the command line.

Command line invocation is a useful feature. Many classes in Sigma include a Unix-style command line interface for testing, or for rare operations that shouldn't clutter up the web-based GUI. Invoking the class with a "-h" parameter to signify "help" will provide a list of allowable options.

```

>$ java -classpath $SIGMA_SRC/build/*:$SIGMA_SRC/lib/*
com.articulate.sigma.KBmanager -h
Sigma Knowledge Engineering Environment
options:
-h - show this help screen
-p - demo Python interface
with no arguments show this help screen and execute a test

```

While loading of the lexicons is straightforward, loading the knowledge bases has several steps. The first is calling the KIF class that parses the SUO-KIF source files, like Merge.kif, which is the original SUMO upper level ontology. The KIF class is responsible for making sure there are no syntax errors in the source, and filling out several indexes. One index lists every term in the knowledge base. Another provides detail about every formula by describing how each term is used. For example

```
(subclass Object Physical)
```

will result in three term pointers. Each will state that the term is found in a "simple" formula - one that is just a tuple. Then it will list the argument position in which it is found. The relation appears in argument "0". So, `Object` will be listed as appearing in a simple formula at position 1.

The indexing system also reports whether a term is in the premise or conclusion of a rule. In the following

```
(=>
  (and
    (connected ?X ?Y)
    (part ?Y ?Z))
  (connected ?X ?Z))
```

the relation `part` is found in the premise (or "antecedent") of a rule. Each key has a simple textual format as defined in `KIF.createKey()` with a prefix "ant-", "cons-" for antecedent and consequent "arg-" if it's an argument in a simple statement, and a rare case of "stmt" for a complex statement that is not a rule. The prefix is followed by the argument number in the case of a simple statement and then the name of the term. So for the first statement above we'd see

```
arg-0-subclass
arg-1-Object
arg-2-Physical
```

and for the second statement, excluding the logical operators

```
ant-connected
ant-part
cons-connected
```

KIF will primarily fill in two of its member variables during parsing. `formulaMap` will be a map of keys from `createKey()` and the string version of the formula itself as a key. The values will be instances of `Formula`. `Formula` is an important class with many operations. It holds the original `String` version of the `Formula` plus several transformations of the `Formula` that will be explained shortly. It also contains information about the lines in the source file from which it came.

After parsing and indexing, Sigma builds some caches for efficiency. Note that the caches are not strictly necessary - SUMO is defined precisely in logic and its semantics are explicit. Any theorem prover needs only the explicit axioms in SUMO to perform valid inferences on the theory. Any IDE or other program using SUMO likewise just needs to implement its formal semantics. But it's wise to take some of the semantics and treat them as optimizations that can be taken advantage of to improve speed.

SUMO has many *transitive* relations. A transitive relation is one for if a relation holds between terms A and B, and also between B and C, then the relation is also true between A and C. Since SUMO is a large theory, there can be transitive chains, especially using the `subclass` relation, that are as many as 10 steps deep. This can result in considerable inefficiency if we need to answer at runtime whether a given class is a subclass of another. So instead, we create a

table of cached relationship so that if B is a subclass of A, and C is a subclass of B etc all the way to Z then we record that Z, Y, X etc are all subclasses of A, B, C etc. There are also more complicated cases in the if N is an instance M and M is a subclass of L then, N is also an instance of L. Worse yet, an **Attribute** might be a **subAttribute** of another term, which in turn is a **subclass** of another term.

SUMO relies on types for the arguments of relations, as in a modern programming language in which analogously, the parameters on methods have type restrictions. So we must maintain the set of types which are allowed for any given relation by recording not just the explicit type, but all its parents. For example

```
(domain part 1 Object)
(domain part 2 Object)
```

This says that the first and second argument of **part** must be instances of **Object** to be valid. Since **Object** is a subclass of **Physical** and **Entity** we must also allow instances of those classes. In all, KBcache has a number of caches to build, which it does in the method **buildCaches()** and then writes out the cached information to a file

```
buildRelationsSet();
buildTransitiveRelationsSet();
buildParents();
buildChildren();
collectDomains();
buildInstTransRels();
buildDirectInstances();
buildDisjointRelationsMap();
writeCacheFile();
```

Next in the initialization process is performing a set of conversions to formulas to make them ready for inference in standard theorem proving systems. This is controlled by **FormulaPreprocessor.preProcess()**.

2.4 Configuration

Configuration of Sigma is first governed by the environment variable **SIGMA_HOME**, which in the default configuration with point to a directory ".sigmakee" in the user's home directory. Configuration of the Tomcat web service is governed by the **CATALINA_HOME** environment variable. The hierarchical directory structure below shows a default installation.

```
home
  user
    workspace
      sigmakee
      sumo
```



```

        WordNetMappings
    .sigmakee
        KBs
        WordNetMappings
var
    tomcat
        webapps
            sigma

```

The "workspace" directory will hold the Git modules for sigmakee and sumo. When "ant" is used to compile sigma from the Git module /home/user/workspace/sigmakee, it creates a .war file (a sort of zip file) that is then copied to the webapps directory. When Tomcat is started, it automatically expands all the .war files in the webapps directory, creating the sigma subdirectory with all the compiled java .class files and .jsp files. Sigma relies on the SIGMA_HOME environment variable to find the .sigmakee directory. During installation of Sigma, the user will copy SUMO .kif files and SUMO-WordNet mapping files under .sigmakee/KBs where the config.xml file also resides.

The config.xml specifies how to set many other configuration parameters for Sigma, and which .kif files to load. Sigma can have multiple named knowledge bases loaded at one time, each of which is composed of one or more constituent .kif files. However, it's most typical just to have one knowledge base named "SUMO". The configuration file parameter "sumokbname" provides the name of the file that Sigma will look for and the primary knowledge base that is assumed at least to contain the Merge.kif file, which is the uppermost level of SUMO.

Default values for configuration parameters are set in KBmanager.setDefaultAttributes().

The full list of configuration parameters, a typical default value and their purpose is now described:

- **adminBrowserLimit** - value = 200 - Each term that is viewed in the Sigma browser can have very few or very many logical statements in which it appears. The term subclass for example, appears in tens of thousands of statements. To prevent the display of common terms from overwhelming the system (and the user), Sigma limits how many are presented at each time. Administrative users have a larger limit than guest users.
- **baseDir** - value = /home/user/.sigmakee - the base data directory for Sigma, under which the KBs directory is found. Currently, there are no other subdirectories.
- **cache** - value = yes - Whether to cache transitive relationships, using the KBcache class.
- **celtdir** - Obsolete

- **dbUser** - value = sa - The user name for the H2 database that is currently used to store login and contact information for users of an installation of Sigma. Passwords are stored with 128-bit one-way encryption only.
- **editorCommand** - Obsolete
- **englishPCFG** - value = /home/user/Programs/stanford-corenlp-full-2015-12-09 - the directory in which the englishPCFG parser model file for Stanford's CoreNLP system is found. This is used by the class **WSD** in word sense disambiguation
- **graphDir** - value = /var/tomcat/apache-tomcat-8.0.41/webapps/sigma/graph - the subdirectory under which graph images will be placed
- **graphVizDir** - value = /usr/bin - The directory in which the executable for the open source GraphViz graph drawing software resides <https://www.graphviz.org>
- **graphWidth** - value = 600 - When a graph rendered by GraphViz is displayed in the browser, how many pixels wide should it be?
- **holdsPrefix** - value = no - whether to convert predicate variables to strict first order by prefixing with **holds**
- **hostname** - value = localhost - The value of the domain name for the host on which Sigma is run, so it can be used in hyperlinks
- **https** - value = false - Whether to make hyperlinks specify the https protocol. When true this requires that the user has configured the server on which Sigma is running for https, including having a certificate from a certificate authority
- **inferenceEngine** - value = /home/user/Programs/E/PROVER/e_ltb_runner - the location of the theorem prover executable
- **inferenceTestDir** - value = /home/user/.sigmakee/KBs/tests - the location for a suite of inference tests, which are files with a .tq extension.
- **kbDir** - value = /home/user/.sigmakee/KBs - The location where .kif files are read from
- **leoExecutable** - value = /home/user/leo - The location where the LEO-II higher order theorem prover executable is found. Note that since this is experimental, it's not in the Sigma distribution.
- **lineNumberCommand** - Obsolete
- **loadCELT** - value = no - Basically obsolete now that all effort on NLP has shifted to SigmaNLP and the Semantic Rewriting approach

- **loadFresh** - value = false - Whether knowledge bases should be loaded and processed from source each time Sigma and Tomcat are restarted, or if they should be loaded from serialized files if the sources haven't been changed. This should be kept as false unless a developer is working on changing caching and indexing approaches.
- **logDir** - value = /home/user/.sigmakee/logs - Currently unused, this parameter specifies where to put error and warning log files. At the moment, error and warnings are found in the logs/catalina.out file under Tomcat.
- **logLevel** - value = warning - Currently unused.
- **nlpTools** - value = yes - Whether SigmaNLP is installed, and therefore should have a hyperlink to it.
- **overwrite** - value = no - This says whether knowledge base files, when loaded through the Manifest.jsp page, should overwrite files of the same name in the \$SIGMA_HOME/KBs directory
- **port** - value = 8080 - Specifies which port Tomcat should run under for Sigma requests. If you're running https, you'll likely use 8443 instead. This number gets put in the hyperlinks in the browser.
- **prolog** - Obsolete
- **semRewrite** - value = /home/user/workspace/sumo/WordNetMappings/SemRewrite.txt
This isn't used by SigmaKEE but is used by SigmaNLP. It probably should be in a separate configuration file.
- **showcached** - value = no - Whether to show cached statement in the browser, or just show statements that are loaded from .kif files
- **sumokbname** - value = SUMO - The name of the SUMO knowledge base.
- **systemsDir** - value = /home/user - Where all the TPTP theorem provers are installed in a local installation of TPTP. This is not exactly obsolete, but hasn't been tested in many years, and most users will not be concerned with it.
- **testOutputDir** - value = /var/tomcat/apache-tomcat-5.5.25/webapps/sigma/tests
- Where to place the results of running inference test questions.
- **TPTPDisplay** - value = no - Whether to display axioms in TPTP format. This is a bit obsolete as it's now possible to do this from the "Formal Language" menu in the Sigma term browser.
- **tptpHomeDir** - value = /home/user

```

(instance part TransitiveRelation)
(<=>
  (instance ?REL TransitiveRelation)
  (forall (?INST1 ?INST2 ?INST3)
    (=>
      (and
        (?REL ?INST1 ?INST2)
        (?REL ?INST2 ?INST3))
      (?REL ?INST1 ?INST3))))

```

Figure 2.3: Example axioms showing predicate variables

- **TPTP** - value = yes - This is where the TPTP framework itself, rather than the theorem provers it communicates with, are found. Most users will not be concerned with this.
- **typePrefix** - value = yes - Whether to add sortal preconditions to axioms as described on page 24
- **userBrowserLimit** - value = 25 - Similar to baseDir above, but a lower limit set for guest users

2.5 Conversion of SUO-KIF to First-Order Logic

Integration of the standard TPTPWorld suite of theorem provers with Sigma [Trac et al., 2008] and the E theorem prover [Schulz, 2002] required processing SUMO into a more strictly first-order syntax. We now discuss the transformations that have been needed.

Although not a required part of the syntax of SUO-KIF, in SUMO, by convention, relations are written with an initial lowercase character, and functions, non-relational instances and classes are written with initial capital letters.

Predicate Variables

First-order provers do not typically support variables in the predicate position. Our first approach was to add a "dummy" predicate to all clauses other than those with logical operators. For example, the axioms in Figure 2.5, which have the variable ?REL in the predicate position, become the axioms in Figure 2.5. This transformation is performed in `FormulaPreprocessor.preProcessRecurse()`.

This transformation however resulted in poor performance for theorem provers that give special indexing priority to the predicate when searching the proof space.

```
(holds instance part TransitiveRelation)
(<=>
  (holds instance ?REL TransitiveRelation)
  (forall (?INST1 ?INST2 ?INST3)
    (=>
      (and
        (holds ?REL ?INST1 ?INST2)
        (holds ?REL ?INST2 ?INST3))
      (holds ?REL ?INST1 ?INST3))))
```

Figure 2.4: Example axioms with holds prefixing

```
(=>
  (holds inverse ?REL1 ?REL2)
  (forall (?INST1 ?INST2)
    (<=>
      (holds ?REL1 ?INST1 ?INST2)
      (holds ?REL2 ?INST2 ?INST1))))
```

Figure 2.5: "holds" insertion

```
(=>
  (holds_3__ inverse ?REL1 ?REL2)
  (forall (?INST1 ?INST2)
    (<=>
      (holds_3__ ?REL1 ?INST1 ?INST2)
      (holds_3__ ?REL2 ?INST2 ?INST1))))
```

Figure 2.6: Holds prefix with arity

```
(=>
  (and
    (part ?INST1 ?INST2)
    (part ?INST2 ?INST3))
  (part ?INST1 ?INST3))
```

Figure 2.7: Example axiom with predicate instantiation

An additional issue is that while KIF-Vampire was customized to support variable-arity predicates, and reuse of names for both predicates and functions, many theorem provers, such as those in the TPTPWorld suite, do not support that flexibility. Translation required creating new predicates for every arity, and a separate set for functions, which are called `holds_X__` and `apply_X__`, respectively, where X is the arity. This transformation has an added benefit of improving performance for those provers which index clauses primarily on the predicate name. The above axiom then becomes as shown in Figure 2.5.

Another approach is to instantiate every predicate variable with all possible values for predicates in the knowledge base that meet the type restrictions that may be implied by the axiom. The rule in Figure 2.5 above will be duplicated with the variable `?REL` being instantiated with every `TransitiveRelation` as in Figure 2.5

If there are ten such relations, there will be 10 copies of the axiom, each with different values for `?REL`. This results in an automated expansion of the number of axioms, but does give good performance. One limitation however is that the semantics of predicate variables is thereby limited to the set of predicates existing in the knowledge base, rather than ranging over all possible predicates.

This transformation is called in `FormulaPreprocessor.replacePredVarsAndRowVars()` and `PredVarInst.instantiatePredVars()` does most of the actual work of the transformation.

```
(=>
  (and
    (subrelation ?REL1 ?REL2)
    (?REL1 @ROW))
  (?REL2 @ROW))
```

Figure 2.8: Axiom with row variables

```
(=>
  (and
    (subrelation ?REL1 ?REL2)
    (?REL1 ?ARG1))
  (?REL2 ?ARG1))

(=>
  (and
    (subrelation ?REL1 ?REL2)
    (?REL1 ?ARG1 ?ARG2))
  (?REL2 ?ARG1 ?ARG2))
```

Figure 2.9: Row variable expansion

Row Variables

Row variables allow us to reference predicates where the number of arguments is not known. While the unbounded implementation of the existence of row variables would make SUO-KIF technically an "infinitary logic" [Hayes and Menzel, 2001], with associated issues in efficient implementation, a bounded interpretation, as described now, does keep SUO-KIF out of this problem.

Sigma treats row variables as "macros", which get expanded automatically so the one axiom in Figure 2.5 becomes several axioms, as shown in Figure 2.5. For brevity only expansion to two variables is shown, but the expansion algorithm continues up to the maximum arity currently allowed of 7, when appropriate. Note that in axioms such as this, which also require predicate variable instantiation, we must restrain the expansion to only those arities which are compatible with the instantiated predicates. For example, `located` is a `subrelation` of `partlyLocated` and both have arity 2. So, `@ROW` will only be expanded to the case of two variables. In the few cases where axioms have two row variables, this can result in 49 new axioms. This work happens in

`RowVars.expandRowVars()`.

Quoting

The original version of KIF had an explicit single quote for denoting uninterpreted structures that were essentially terms. This was used to state complex expressions which could be read by humans, without incurring the computational cost of becoming higher-order logic. For example `(believes Mary (likes John Sue))` is a higher-order expression, because the second argument to `believes` is not a term. `(believes Mary (likes John Sue))` is first-order in the original KIF because the single quote character converts the following list into a term. This however is not strictly necessary since a reasoning system can apply a quote automatically when needed by looking at the form of the arguments, or the domain statements which define the argument types of the predicate. SUO-KIF allows the unquoted expression and leaves it to a reasoning system how it wishes to handle it. If a higher-order interpretation is possible, then that is allowed. If not, then the reasoning system is responsible for quoting any argument to a relation which is not a term. Sigma employs the latter approach when sending statements to its suite of first-order logic reasoners.

Quoting removes most of the semantics of higher-order statements, including the semantics of logical operators, but does at least allow for unification, thereby giving the appearance of higher-order reasoning in very limited situations.

For example,

`(believes John (likes Mary Jeff))`

becomes

`(believe John '(likes Mary Jeff))`

This allows KIF-Vampire to perform very simple queries on higher-order statements, such as

`(believes John '(likes Mary ?X))`

and get the correct answer of `Jeff`. However, logical symbols in the embedded formulas lose their meaning, so if

`(believes John '(and (likes Mary Jeff) (likes Bill Sue)))`

is asserted, the same query will fail, as the `and` does not have its conventional meaning, and the two lists will not unify.

A better way to address these constructs is to treat them with their full higher-order semantics. For this we translate SUO-KIF to the THF language in the class `THF` and send them to the LEO-II prover for inference.

Sortals

Provers such as KIF-Vampire are unsorted. That is, variables may range over any type. However, SUMO specifies the types of arguments required for each


```
(=>
  (and
    (instance ?TRANSFER Transfer)
    (agent ?TRANSFER ?AGENT)
    (patient ?TRANSFER ?PATIENT))
  (not
    (equal ?AGENT ?PATIENT)))
```

Figure 2.10: Example axiom

```
(=>
  (and
    (instance ?AGENT Agent)
    (instance ?PATIENT Object))
  (=>
    (and
      (instance ?TRANSFER Transfer)
      (agent ?TRANSFER ?AGENT)
      (patient ?TRANSFER ?PATIENT))
    (not
      (equal ?AGENT ?PATIENT))))
```

Figure 2.11: Sortal prefixing

predicate. When run in an unsorted prover, these specifications have the unintended effect of generating contradictions. Because variables can be of any type, they may, during search, be bound to a type that is incompatible with the restrictions on a particular predicate's argument types that are also part of the search. The axiom that specifies the argument type restriction for that predicate may then contradict that variable binding. In addition, by allowing variables to be any type, search may include finding bindings for variables that cannot be part of the eventual successful solution, so there is an efficiency cost, as well as a problem for finding an accurate proof.

We should note that there is an efficiency cost with using sortal prefixes also, since they increase the number of literals that must be proved in order to derive each conclusion. We would expect the use of sortal prefixes to improve correctness, but at the cost of speed (and some space). Perhaps the use of sortals has not provided any obvious benefit so far only because we have not allowed each test to run for a long enough time.

To solve this problem of having search unconstrained by argument types,

```

(=>
  (and
    (instance ?AGENT Agent)
    (instance ?TRANSFER Instance)
    (instance ?TRANSFER Process)
    (instance ?PATIENT Object))
  (=>
    (and
      (instance ?TRANSFER Transfer)
      (agent ?TRANSFER ?AGENT)
      (patient ?TRANSFER ?PATIENT))
    (not
      (equal ?AGENT ?PATIENT))))

```

Figure 2.12: Naive sortal prefixing

we generate additional preconditions for each rule in the ontology, which then limits every rule to being considered only if type requirements have been met. For example, Figure 2.5 is transformed into Figure 2.5.

Note that a nave implementation of this approach would be to state the version in Figure 2.5 but since `?TRANSFER` is already further constrained by the first clause of the original rule, those additional clauses are not necessary.

This processing is handled in `FormulaPreprocessor.addTypeRestrictions()`.

After these pre-processing steps are performed to transform row variables, instantiate predicate variables and add sortal prefixes, Sigma then can translate the formulas to the TPTP language for use in first-order theorem proving. This is handled in `SUMOKBtoTPTPKB` which then calls on `SUM0formulaToTPTPformula`. The results of this processing are stored in `Formula.theTptpFormulas`.

2.6 Normalization Algorithm

Normalization is a process of reducing the complexity of logical formulas by removing logical symbols which are not strictly needed. This makes the logic harder to read, so it's useful to be able to author knowledge in one form but then reduce it to a simpler form for automated processing. All modern theorem provers work with a so-called normal form, which removes quantification and implication symbols. We use Russell and Norvig's algorithm [Russel and Norvig, 2009].

Normalization is handled in the class `Clausifer`. Sigma makes limited use of the class, primarily just for a rare, expensive but more precise test for equality between formulas than just looking at textual matches. In the descriptions below we'll provide TPTP syntax on the left and SUO-KIF syntax on the right. Each

step corresponds to a method in the code. The top level call to `clausify()` is

```
public Formula clausify() {
    thisFormula = equivalencesOut();
    thisFormula = implicationsOut();
    thisFormula = negationsIn();
    thisFormula = renameVariables();
    thisFormula = existentialsOut();
    thisFormula = universalsOut();
    thisFormula = disjunctionsIn();
    thisFormula = standardizeApart();
    return thisFormula;
}
```

Remove implications and equivalences

$a \Leftrightarrow b$ is the same as $(\Leftrightarrow A B)$ is the same as
 $(a \Rightarrow b) \wedge (b \Rightarrow a)$ (and $(\Rightarrow A B) (\Rightarrow B A)$)
 $a \Rightarrow b$ is the same as $\neg a \vee b$ $(\Rightarrow A B)$ is the same as $(\text{or } (\text{not } A) B)$

Move negation inwards

$\neg(p \vee q)$ becomes $\neg p \wedge \neg q$ (not (or P Q)) becomes
 (and (not P) (not Q))
 $\neg(p \wedge q)$ becomes $\neg p \vee \neg q$ (not (and P Q)) becomes
 (or (not P) (not Q))
 $\neg![X]:p$ becomes $?[X]:\neg p$ (not (forall (?X) P)) becomes
 (exists (?X) (not P))
 $\neg?[X]:p$ becomes $![X]:\neg p$ (not (exists (?X) P)) becomes
 (forall (?X) (not P))
 $\neg\neg p$ becomes p (not (not P)) becomes P

Standardize variables

The scope of variables in quantifiers is local to the quantifier, but to avoid confusion, variables of the same name but different scope in a formula are renamed.

$![X]:p(X) \vee ?[X]:q(X)$ becomes (or
 $![X]:p(X) \vee ?[Y]:q(Y)$ (forall (?X)
 (p ?X))
 (exists (?X)
 (q ?X)))
 becomes
 (or
 (forall (?X)
 (p ?X))

```
(exists (?Y)
  (q ?Y)))
```

Move quantifiers left

$p \mid ! [X] q(X)$ becomes $(\text{or } P \text{ (forall } (?X) (q ?X)))$ becomes
 $! [X] p \mid q(X) \text{ (forall } (?X) (\text{or } P (q ?X)))$

Skolemization

Create a unique function in place of every existentially quantified variable. Include every universally quantified variable that is in scope, inside the function.

```
! [X]:p(X) => (forall (?X)
  (?[Y]:h(Y) \& a(X,Y)) (=>
    (p ?X)
    (exists (?Y)
      (and
        (h ?Y)
        (a ?X ?Y))))))
```

becomes

```
! [X]:p(X) => (forall (?X)
  (h(skf(X)) \& a(X,skf(X))) (=>
    (p ?X)
    (and
      (h (skf ?X))
      (a ?X (skf ?X)))))
```

Distribute and over or

$(a \ \& \ b) \mid c$ becomes $(\text{or } (\text{and } a \ b) \ c)$ becomes
 $(a \mid c) \ \& \ (b \mid c) \text{ (and } (\text{or } a \ c) (\text{or } b \ c))$

Flatten nested conjunctions and disjunctions

$(a \mid b) \mid c$ becomes $(\text{or } (\text{or } a \ b) \ c)$ becomes
 $a \mid b \mid c \text{ (or } a \ b \ c)$

2.7 Conversion to TPTP

While the conversions described above result in an essentially first-order form, there are several aspects that are beyond the "traditional human-readable" format of the TPTP language, as parsed by many current provers. The TPTP language uses Prolog-like user terms and atoms, uses infix notation for binary

operators, has a separate namespace for operators, and provides a separate namespace for defined functors and predicates. Additionally the TPTP language does not support arbitrary lists. These differences are dealt with in the translation to TPTP format as follows.

A recursive algorithm is used to convert from the SUO-KIF prefix form for binary operators, stacking the translated form of an operator when found at the start of a formula, copying it off the top of the stack for insertion between operand formulae, and popping it off the stack at the end of the formula. As user terms and atoms are encountered they are translated to Prologs prefix form, with functors and predicates first letters being set to lowercase, and variables first letters being set to upper case. All hyphens in user terms are translated to underscores. Functions, recognized in SUO-KIF by the "Fn" suffix, are translated to corresponding equivalents from the TPTP language, starting with a "\$". A minor translation of double quoted strings is performed, replacing non-printable characters - carriage return, new line, tab, and formfeed - by spaces.

Truly higher-order constructs are dealt with by losing most of their semantics by conversion to uninterpreted lists. This translated form is not directly usable in the TPTP format, as there is no support for arbitrary lists. The current solution is to lose even more of the semantics, by single quoting such expressions, thus treating them as constants. In this way the possibility of unification over the list elements is lost - only unification of the whole is possible. Part of the reason for taking this simplistic approach was to permit consistent translation of operators, and because they have a separate namespace in the TPTP language they cannot be treated as terms in a list. The list solution can be implemented in the translation to TPTP format by retaining the SUO-KIF forms of operators (which look like TPTP constants), adding the `_X_` suffix to the predicate list to form unique predicates for different arities at the atom level, and providing `listf_X_` functors for nested lists.

Chapter 3

Using Sigma

3.1 JSP Interface

Sigma has a number of functions controlled through its JSP-based interface. Sigma runs under Tomcat, which is a web server that allows Java code to be embedded in HTML pages and run on the server, as opposed to JavaScript, which is run on the client.

- `AddConstituent.jsp` - adds a constituent SUO-KIF file to a knowledge base. Accessible only to admin users. It just responds to a command from another page and has no UI of its own. Redirects back to `KBs.jsp` when done.
- `AllPictures.jsp` - shows all the pictures linked to a term at once. Accessible to all users.
- `ApproveUser.jsp` - handles approving a new user. Accessible only to admin users. Redirects to `KB.jsp` once acknowledge.
- `AskTell.jsp` - interface to the local theorem provers. Accessible only to admin users. This function has not been well maintained and the interfaces to `SystemOnTPTP` and `LEO-II` may be out of date.
- `BrowseBody.jsp` - shows terms, axioms, lexicon links, etc
- `BrowseExtra.jsp` - includes `Prelude.jsp` and `Postlude.jsp`
- `BrowseHeader.jsp` - primary ontology browser controls including term and word search, as well as the menu for selecting natural language and formal language
- `Browse.jsp` - top level browsing JSP that includes `Prelude.jsp`, `BrowseHeader.jsp`, `BrowseBody.jsp` and `Postlude.jsp`. Really just a shell for the included JSPs

- CCheck.jsp - interface to `KBmanager.initiateCCheck()` that initiates consistency checking of a KB. Handles selection of which theorem prover to use, and several parameters. Includes Prelude.jsp and Postlude.jsp
- CELT.jsp - Obsolete. Handles invocation of the Controlled English to Logic Translation system, which is now superceded by the Semantic Rewriting approach in the sigmanlp project.
- CreateUser.jsp - handles a request from a user to create an account. Creates a guest account and sends mail to the moderator for approval. Has no UI.
- Diag.jsp - Interface to run tests in Diagnostics.java. Depends on Prelude.jsp and Postlude.jsp. Accessible for "admin" and "user" but not unregistered guest users
- EditFile.jsp - Obsolete
- EditStmt.jsp - Obsolete
- Graph.jsp - Create graphs of binary relationships as a graphical view and as indented text. Relies on Prelude.jsp and Postlude.jsp
- InferenceTestSuite.jsp - run a suite of inference tests with the prover and parameters selected. Depends on Prelude.jsp and Postlude.jsp. Accessible only for admin users.
- init.jsp - A status page with periodic automatic refresh to catch requests to Sigma during the process of initialization.
- InstFiller.jsp - Page to allow simple editing of ground formulas. Deprecated. Accessible only for admin users. Depends on Prelude.jsp and Postlude.jsp.
- Intersect.jsp - Find appearances of two or more terms in the same axiom. Depends on Prelude.jsp and Postlude.jsp.
- KBs.jsp - Main entry point for creating or selecting a knowledge base. Some functions are available for "admin" or "user" roles but not "guest". Depends on Prelude.jsp and Postlude.jsp.
- login.html - login page for Sigma that handles existing accounts and new account registration. Existing accounts are dispatched to login.jsp and registrations are dispatched to Registration.jsp
- login.jsp - handle the login process by using PasswordService.jsp to validate a login against an H2 database of account info. This page has no UI.

- Manifest.jsp - UI for handling loading and saving KB constituents including saving in various exportable formats, such as OWL. Depends on Prelude.jsp and Postlude.jsp. Most functionality is limited to admin users.
- Mapping.jsp - Use some simple string matching approaches to suggest equivalences between terms in two files. Depends on Prelude.jsp and Postlude.jsp. Access limited to admin users.
- MiscUtilities.jsp - Depends on Prelude.jsp and Postlude.jsp. Most functionality is limited to admin users. Little utilities to generate dot graph files and OWL versions of the Open Multilingual Wordnet content.
- ModeratorApproval.jsp - Functionality is limited to admin users. Dispatches to ApproveUser.jsp when user clicks "ok".
- OMW.jsp - Display results from the many languages in Open Multilingual Wordnet linked to WordNet synsets. Depends on Prelude.jsp and Postlude.jsp.
- OWL.jsp - Display all the axioms for a given term in OWL format, if expressible in that language.
- Postlude.jsp - Encapsulates footing information displayed on most pages.
- Prelude.jsp - Encapsulates header information displayed on most pages.
- ProcessFile.jsp - No UI. Called from MiscUtilities.jsp to generate KIF from other formats. Calls DocGen.dataFileToKifFile() to do the real work. Depends on Prelude.jsp and Postlude.jsp. Access limited to admin users.
- Properties.jsp - An interface for setting many of Sigma's parameters. Some of these are also accessible from the browsing interface, such as the language for the natural language paraphrases. Depends on Prelude.jsp and Postlude.jsp. Access limited to admin users.
- Register.jsp - An interface to allow people to submit a request for registration and privileges beyond that of unregistered guest users. Calls on CreateUser.jsp
- Save.jsp - not used
- SimpleBrowseBody.jsp - Analogue to BrowseBody.jsp but showing only simple axioms in a simple and non-technical language
- SimpleBrowseHeader.jsp - Analogue to BrowseHeader.jsp for simple axioms
- SimpleBrowse.jsp - Analogue to Browse.jsp for simple axioms

- `SystemOnTPTP.jsp` - Interface to the SystemOnTPTP system hosted at U. Miami that collects dozens of theorem provers with a common programmatic interface. This code has not been maintained so use at your own risk. Depends on `Prelude.jsp` and `Postlude.jsp`. Access limited to admin users.
- `TreeView.jsp` - A simple tree view of the taxonomic structure of the ontology with collapsable nodes. Uses the `SimpleBrowse.jsp` code to display axioms for any node in the taxonomy.
- `WN Diag.jsp` - Diagnostics for WordNet. Depends on `Prelude.jsp` and `Postlude.jsp`. Accessible for "admin" and "user" but not unregistered guest users.
- `WordNet.jsp` - show synsets and SUMO-WordNet mappings for a word. Depends on `Prelude.jsp` and `Postlude.jsp`.
- `WordSenseFile.jsp` - show word sense disambiguation and sentiment analysis results for a file of text. Calls `WordNet.wn.sumoFileDisplay()` for the real work. Depends on `Prelude.jsp` and `Postlude.jsp`.
- `WordSense.jsp` - show word sense disambiguation and sentiment analysis results for a sentence. Calls `WordNet.wn.sumoSentenceDisplay()` for the real work. Depends on `Prelude.jsp` and `Postlude.jsp`.

3.2 Language Generation

SUMO proper has a significant set of manually created language display templates that allow terms and definitions to be paraphrased in various natural languages, including those with non-western character sets. This approach was first developed in [Sevcenko, 2003]. These templates now include Arabic, French, English, Czech, Tagalog, German, Italian, Hindi, Romanian and Chinese (traditional and simplified characters). Automatically generated natural language paraphrases can be seen in the rightmost column of the screen display given as Figure 3.1.

This functionality is implemented in the `com.articulate.sigma.nlg` package, primarily in the `LanguageFormatter` class. There is an expanded approach for NLG in the remaining classes in the package, but it is not fully implemented so those classes should be ignored for now.

Although presentation of terms is straightforward, presentation of statements is more complicated, and roughly patterned after C-language printf statements. For example the relation part, which states that one object is a part of another, has the corresponding language generation template "%1 is %n a part of %2". The first argument to the logical relation is substituted for %1 etc. Note that this substitution is recursive, so that complex statements with nested formulae can be translated effectively. The %n signifies that if the statement is

The screenshot shows the Sigma browser interface. At the top, there's a navigation bar with 'Home' and 'Graph' links. Below it, the 'KB' is set to 'SUMO' and the 'Language' is 'EnglishLanguage'. The main content area is titled 'Walking (walking)' and includes two images: a person walking and a robot walking. To the right of the images is a list of synonyms: Rollerblade, stroll, amble, ambulate, ambulation, angry, walk, backpack, break, lumble, canter, career, circumambulate, clamber, climb, climb up, clomp, clump, cock, coogle, constitutional, constitutionalize, counter-march, crab, creep, crawl, dash, hobnob, dodder, dogtrot, drag, dressage, drift, err, escalate, exhibit, falter, fast break, file, file in, file out, fire walking, founce, founder, foot, hater, footslog, footstep, forage, gait, gallop,...

Below the images, there are two sections: 'appearance as argument number 1' and 'appearance as argument number 2'. The first section lists various external images and their URLs. The second section lists various external images and their URLs. At the bottom, there is a 'consequence' section with a list of logical rules.

KB Term: Walking **Show**

English Word: **Noun** **Show**

Walking (walking)

Rollerblade, stroll, amble, ambulate, ambulation, angry, walk, backpack, break, lumble, canter, career, circumambulate, clamber, climb, climb up, clomp, clump, cock, coogle, constitutional, constitutionalize, counter-march, crab, creep, crawl, dash, hobnob, dodder, dogtrot, drag, dressage, drift, err, escalate, exhibit, falter, fast break, file, file in, file out, fire walking, founce, founder, foot, hater, footslog, footstep, forage, gait, gallop,...

appearance as argument number 1

(documentation Walking EnglishLanguage "ambulating relatively slowly, i.e. moving in such a way that at least one foot is always in contact with the ground.") Merge.kif 8825-8826

(externalImage Walking "http://upload.wikimedia.org/wikipedia/commons/0/0f/Robotpeintre.gif") pictureList.kif 3030-3030

(externalImage Walking "http://upload.wikimedia.org/wikipedia/commons/6/6f/Walk-Cycle.gif") pictureList.kif 3276-3276

(externalImage Walking "http://upload.wikimedia.org/wikipedia/commons/d/d2/Marcheur_en_comp%C3%A9ton.jpg") pictureList.kif 3277-3277

(subclass Walking Ambulating) Merge.kif 8824-8824

Walking is a subclass of ambulating

appearance as argument number 2

(partition Ambulating Walking Running) Merge.kif 8819-8819

(subclass Wading Walking) Mid-level: ontology.kif 236-236

(termFormat EnglishLanguage Walking "walking") english_format.kif 792-792

Ambulating is exhaustively partitioned into walking and running

Wading is a subclass of walking

term format english language, walking and "walking"

antecedent

(=>

(and

(instance ?WALK Walking)

(instance ?RUN Running)

(agent ?WALK ?AGENT)

(agent ?RUN ?AGENT)

(holdsDuring

(WhenEn ?WALK

(measure ?AGENT

(SpeedEn ?LENGTH1 ?TIME)))

(holdsDuring

(WhenEn ?RUN

(measure ?AGENT

(SpeedEn ?LENGTH2 ?TIME)))

(greaterThan ?LENGTH2 ?LENGTH1))

Merge.kif 8833-8841

- If a process is an instance of walking and process is an instance of running and an agent is an agent of process and agent is an agent of process and the measure of agent is a length measure per a time duration holds during the time of existence of process and the measure of agent is length measure per time duration holds during the time of existence of process
- then length measure is greater than length measure

consequence

Find: **Previous** **Next** **Highlight all** ☐ Match case **Reached end of page, continued from top**

Done

Figure 3.1: Sigma browser screen

negated, that the negation operator for the appropriate human language should be inserted in that position.

The predicate format associates a concept (either a relation or a function) with a string. `format` takes three arguments: the name or abbreviation of a natural language, the relation name and the format string. When there is a need to render a concept in natural language, the associated string is used. The string contains a natural language description of the concept and special tags which are interpreted with the browser.

Take for example that we have the SUO-KIF statement that

```
(authors Dickens OliverTwistBook)
```

We have the following statements that have been coded to support the paraphrasing of statements with the authors relation.

```
(format EnglishLanguage authors
  "%1 is %n the %&%author of %2")
(format it authors "%1 l' %&%autore di %2")
```

Terms are also given language-specific strings, when appropriate

```
(termFormat EnglishLanguage OliverTwistBook "Oliver Twist")
```

If a Sigma user has loaded this information in a knowledge base, and English is selected as the presentation, the user will see "Dickens is the author of Oliver Twist." next to the SUO-KIF statement. If Italian is selected, the paraphrase will be "Dickens l'autore di Oliver Twist". As mentioned above, the refers to the word for negation in the given language, and is inserted if the formula is negated. For example,

```
(not (authors RobinCook WarAndPeace))
```

is rendered as "Robin Cook is not the author of War and Peace." The full description of tags that may be included in the format statements is as follows:

- `&%token` specifies a token that will be made into a hypertext link to the concept being visualized.
- `%1, %2, ...` this tag will be substituted with a natural language representation of the concepts respective argument.
- `%ntext` will be replaced either with an empty string, if a predicate is being rendered as positive, or "text" otherwise; the `%n` tag can be used as a shortcut for `%nnot`. `%ptext` will be replaced with "text" for positive rendering and with an empty string for negative rendering.

```
(=>
  (and
    (instance ?AMBULATE Ambulating)
    (agent ?AMBULATE ?AGENT))
  (attribute ?AGENT Standing))
```

Figure 3.2: An axiom for *Ambulating*

- `%*range[delim]` will be replaced with a list of natural-language representations of a subset of arguments; range specifies which arguments will be included; it is a comma separated list of numbers or ranges, for example, "1-4,6" denotes the first, second, third, fourth and sixth argument; the `delim` parameter specifies the delimiter which will be used to separate representations of arguments; both range and `[delim]` may be omitted, in which case range defaults to all arguments, and `[delim]` defaults to a single space.
- `%%` will be replaced with a single percent character.

This is a simple approach that works reasonably well for English, but which does not work well for languages with more types of grammatical agreement.

More recently, we have developed an approach to improve the translation of axioms involving variables. Variable names can be any string, but as they are often given the form of English words, it was confusing to present the names unaltered in presentations other than English. The solution was to collect the types of variables and use that type as its name. Take for example the axiom in Figure 3.2.

Sigma will generate the following English paraphrase: "If a process is an instance of ambulating, and an agent is an agent of the process, then standing is an attribute of the agent."

Sigma checks the argument type restrictions of the relations to determine the type of the variables. For example, the first argument of the agent relation is an instance of a **Process** and the second argument is an instance of an **Agent**. Sigma uses the `termFormat` expression for **Process** to generate the string `process`". If instead the user had asked for an Italian paraphrase the string `processo`" would be the result. This is preferable to the old version of the natural language paraphrase in which the user would have seen the variable name `?AMBULATE` even if a non-English paraphrase was requested.

3.3 Browsing and Display

Sigma was originally just a display tool. Its original, and still most heavily used function, is for creating hyperlinked sets of formatted axioms that all contain

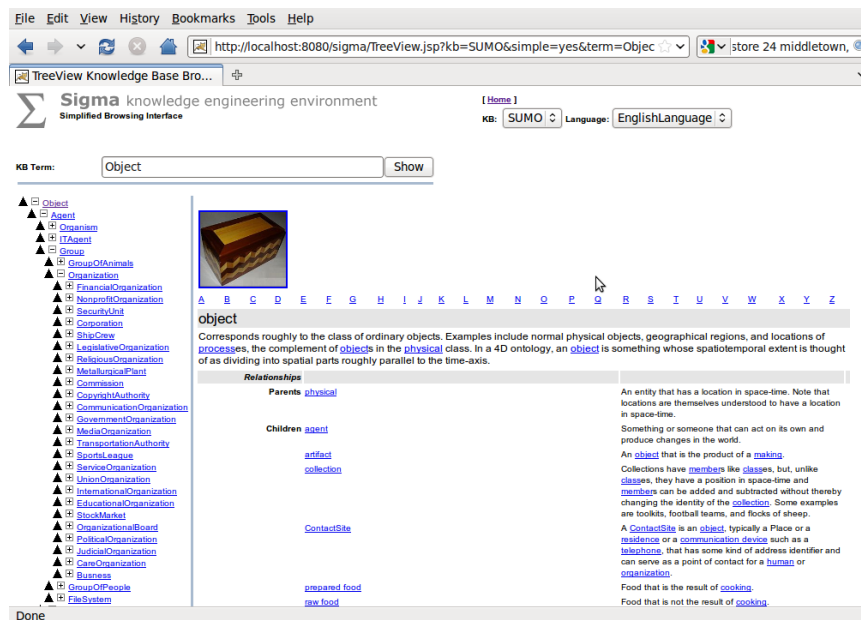


Figure 3.3: Simplified browser view

a particular term (Figure 3.1). Clicking on a term in turn gives a hyperlinked display of all the axioms that contain the new term. Next to each axiom is given the file and line(s) where the axiom was written. Also shown is an automatically generated natural language paraphrase of each axiom.

A simplified browser view (Figure 3.3) that may be more appropriate for users who are transitioning from use of frame and description logic languages. It gives prominence to a tree view of the subclass hierarchy and presents binary relations in a simple tabular format, relegating rules to an area lower in the browser pane, and rendering them in the natural language paraphrase form only.

SUMO has several hierarchies that can be used to organize and display the theory. These include hierarchies of physical parts, relations, attributes, processes and others. As such, the tree browser allows the user to select any transitive binary relation as the link by which the hierarchy display is created.

Images provide an informal visual representation of as many of the concepts in SUMO as possible. Some 12,000 links were made by hand to public domain icons and images in Wikipedia. Some 900 images linked to Wikipedia in Princeton's ImageNet [Deng et al., 2009] corpus were also imported.

3.4 Analysis and Debugging

Sigma includes a number of specialized and general tools for ensuring ontology quality. These are contained in the `Diagnostics` class. The ultimate tool for quality checking on a formal ontology is formal reasoning. However, in expressive ontologies, such as SUMO, we can generally not expect that all contradictions can be detected with theorem provers or that consistency can be formally proved (note, for example, that Peano arithmetic can be formalized in SUO-KIF). Sigma therefore provides different tools for quality checking, combining exhaustive and terminating special purpose tests with incomplete and generally non-terminating general purpose testing based on theorem proving or model finding.

There are two special case tests for errors that must be corrected. We test for terms without a root in the subclass hierarchy at the term `Entity`, which is the topmost term in SUMO (in `Diagnostics.termsNotBelowEntity()`). This commonly results from either omitting a subclass or instance statement when defining a new term, or by misspelling the name of the intended parent term. The second special case test is for where a term has parents that are defined to be disjoint (in `Diagnostics.childrenOfDisjointParents()`). In a large theory like SUMO, it can be easy to lose track of this case, especially when the ultimate conflict may be between terms that are many levels up in the subclass hierarchy.

There are also a number of tests for cases that are indicative of a problem, yet not strictly an error that would result in a logical contradiction. The first of these is for terms lacking documentation (in `Diagnostics.termsWithoutDoc()`). In theories under construction, theories that are the results of importing and merging another ontology, or simply for large lists of domain instances, like city names, it may be reasonable, temporary, or expected for such terms to lack documentation. But this does often reflect an outright error, where a term name was simply misspelled in the documentation definition, or in some other axiom. The test `Diagnostics.termsWithMultipleDoc()` may also be indicative of a problem, as can relations without any natural language `format` expression (checked in `Diagnostics.relationsWithoutFormat()`).

We test for cases where terms do not appear in any rules (in `Diagnostics.termsWithoutRules()`). This again is common in collections of instance-level facts, but undesirable for many classes or relations, where it should be possible to define precisely the intended meaning of the term with a small number of formal rules, as well as statements like class membership.

Rules can be syntactically valid with quantified variables that appear in a quantifier list but aren't used in the scope of the quantifier. These can be found with `Diagnostics.quantifierNotInBody()` and are often indicative of a typo in the variable name.

A more complicated issue is where members (instances) of a parent class that are not also members of one of the subclasses that constitute the exhaustive decomposition of the parent class (stated with

the relation `exhaustiveDecomposition`). This problem is found with `Diagnostics.membersNotInAnyPartitionClass()`.

Because knowledge bases are often composed from SUMO’s general and domain specific component ontologies, it is desirable to limit dependencies among the files as much as possible. For that reason we include a tool to specify dependencies between pairs of files (in `Diagnostics.printTermDependency()`). It is typically most desirable at least to ensure that dependencies are only from one file to another, and not between both files. All domain files will of course depend at least upon SUMO proper, since they form a single integrated theory that is decomposed into separate files for convenience and efficiency of inference.

A further test exploits the SUMO-WordNet mappings. They offer the opportunity to find problems exposed by differences in the two products. The two hierarchies should not necessarily be isomorphic, and therefore respective differences do not necessarily mark an error, but are worth examination when they are found.

In the diagnostics provided for the SUMO-WordNet mappings (found in `WNdiagnostics`), Sigma finds WordNet synsets without mapped formal terms (in `WNdiagnostics.synsetsWithoutTerms()`) and those for which a formal term is provided, but is not found in the current loaded knowledge base (in `WNdiagnostics.synsetsWithoutFoundTerms()`). This helps to find cases where terms have been changed or renamed and the mappings not updated. There is also a reverse case of where a SUMO term doesn’t have a synset. Relations in SUMO (including functions) won’t have mappings. There are also terms in SUMO, especially the domain ontologies, which are too specific to have a mapping - only their parent terms in SUMO are mapped and therefore their mapping can be inferred. This case is checked for in `WNdiagnostics.nonRelationTermsWithoutSynsets()`.

Most significant is the taxonomy comparison component. Given that we have terms A and B in SUMO and synsets X and Y in WordNet, if A is mapped to X and B to Y, Sigma checks whether if B is a subclass of A then Y is also a hyponym of X. The reverse case is also checked. This test is found in `WNdiagnostics.nonMatchingTaxonomy()`.

3.5 Inference

The SUMO Inference Engine (SInE) [Hoder and Voronkov, 2011], which selects only the subset of axioms likely to be relevant for a given query. The algorithm is implemented in the SInE class. Results show dramatic improvement on SUMO-based inference over other approaches [Pease and Sutcliffe, 2007]. We should note however that this class is only useful if one wanted to implement a connection to a prover other than Vampire or Eprover, as both of them have incorporated with SInE algorithm since its performance advantage was immediately clear on large knowledge bases.

We have the class `InferenceEngine` to structure interaction with theorem provers. The class `EProver` handles interaction with E. Most users will not

Sigma knowledge engineering environment
SystemOnTPTP Interface

[Home | Graph | Prefs] KB: SUMO Language: ▾

(instance ?X PrimaryColor)

● Local SystemOnTPTP ● Remote SystemOnTPTP System: EP---0.999 ▾

Maximum answers: 2 Query time limit: 300

● TPTP Proof ● IDV-Proof tree ● Hyperlinked KIF

Ask Tell

(SZS Status Theorem)
 Answer 1. [definite] ?X6 = Red

1. (exists (?X6)	[Query]
(instance ?X6 PrimaryColor)	
2. (instance Red PrimaryColor)	[KB]
3. (instance Red PrimaryColor)	2
4. (not	1
(instance ?X6 PrimaryColor)	
5. (not	4
(instance Red PrimaryColor)	
6. True	3 5

(SZS Status Theorem)
 Answer 1. [definite] ?X = Blue

1. (instance Blue PrimaryColor)	[KB]
2. (exists (?X1)	[Instantiated
(instance Blue PrimaryColor)	Query]
3. (not	2
(exists (?X1)	
(instance Blue PrimaryColor))	
4. (not	3
(instance Blue PrimaryColor)	
5. (instance Blue PrimaryColor)	1
6. (not	4
(instance Blue PrimaryColor)	
7. True	6 5
8. True	7
9. True	8

Figure 3.4: Proof presentation in Sigma, from [Trac et al., 2008]

need the **Vampire** class, since it only interfaces with KIF-Vampire which is now over a decade old. However, the **SystemOnTPTP** class can be used to interface with the latest Vampire and over 40 other provers, albeit in a somewhat more heavyweight fashion of remote communication to the TPTP servers at U. Miami or by replicating their server installation locally.

In addition to preprocessing, some post-processing is needed for all theorem provers that are used in Sigma. All the TPTP provers that report full proofs, as well as KIF-Vampire, present an ordered list of deductions, where premises are given and then a conclusion. In presenting a proof to the user (Figure 3.4), we would like to avoid showing the same axiom many times if it is used in several proof steps. We therefore assign a numerical index to each axiom, in order of its appearance in the proof. The indexes can then be referenced when they are preconditions to a listed step, making the proof appear more similar to what a logic student will be used to from a standard textbook presentation of a proof.

An answer variable is a binding for a variable in a query that is unbound. In the proof shown in Figure 3.4 ?X is an unbound variable in the query. For TPTP systems that do not report answer variables, or handle more than one answer per query, a more complicated approach is needed. For systems such as EP, that

```
(and
  (instance ?X PrimaryColor)
  (not
    (equals ?X Red)))
```

Figure 3.5: New query, negating the first result in order to find a new answer.

```
(fatherOf John Bill)
(fatherOf John Mark)
(=>
  (fatherOf ?X ?Y)
  (and
    (not
      (exists (?Z)
        (and
          (fatherOf ?X ?Z)
          (not
            (equal ?Y ?Z)))))))
```

Figure 3.6: Example contradictory knowledge base

report proofs but not answer variables, the axioms in the proof are resubmitted to the Metis prover [Hurd, 2003] which does report answer variables. Multiple answers are found by resubmitting the query with a new clause added that excludes previous answers. For the example query shown in Figure 3.4, in order to get the second answer, the new query would become as shown in Figure 3.5.

At the boundary of diagnostics and inference we have the general case of using theorem proving to find contradictions. Because first-order proving is not guaranteed to find all problems that may exist in SUMO, Sigma includes a consistency check function that leads the theorem prover to consider each axiom in a knowledge base. This is an important point because a user may have a knowledge base that is inconsistent, but in practice may make many useful inferences over a long period of time while never having the problem show up in a proof. For example, take the knowledge base shown in Figure 3.5.

While this knowledge base is trivially small, imagine that there are tens or hundreds of thousands of other statements, many of which may involve the predicate symbol `fatherOf`. A complete examination of the proof space is impossible. Imagine that the user poses the query

(fatherOf John Bill)

Given a finite and small amount of time with which to find an answer, the prover may just find and return "yes" after encountering the first assertion. It may not continue the search process to find the contradiction. In fact, given a large and complex enough knowledge base, and a complex enough contradiction, it might not be found for years.

To help guide the search for contradictions, Sigma takes each axiom, which is loaded one by one starting with an empty knowledge base. For each axiom, the prover is asked to compute whether the knowledge base contradicts the axiom, or is redundant with it. If the axiom doesn't create a contradiction, it is asserted to the knowledge base and the next axiom is considered. A contradiction will stop processing, since once a contradiction is found, any further results may be nonsensical (although the answer also may not be nonsensical, as we have explained, so this is a conservative approach).

Once processing finishes, redundancies are collected and reported. At its simplest, a redundancy can be a duplicated statement, and that is clearly an error. Although initially harmless, having the same statement in two places can easily lead to problems as an ontology evolves, as one statement might get changed while a duplicate does not. For example, a developer might forget that a domain ontology file already has a statement (`subclass Table Furniture`) and assert the same statement in a different file.

A more complex case is where one statement is simply deducible from several others. This is often intentional, as knowledge engineers may wish to short-circuit a common chain of reasoning in order to have faster inference. Such a case is even more likely to suffer from the problem of changes not being reflected in the chain of deductions, and the redundant conclusion. For example, see Figure 3.5.

Similar to the CASC competition, but on a much smaller scale, Sigma has the capability to run a series of SUMO-based tests for any theorem prover it supports, reporting success or failure and the time taken on each test. This is handled in the `CCheck` and `CCheckManager` classes.

3.6 Proof Presentation

Sigma includes several options for proof presentation. Despite the fact that most textbooks present proofs as linear structures, proofs are trees. The same sub-proof may be used several times in reaching different intermediate conclusions. In order to ease the understanding of Sigma proofs for humans, Sigma's proof linearization mechanism therefore integrates tools that detect and eliminate repetitions of identical sub-proofs.

When using the TPTP system integrated into Sigma, the user also has the option of a graphical proof presentation with the IDV component [Trac et al., 2007].

Sigma has options for how it presents multiple proofs. Often, the same proof, with the same proof steps, can be reached via a different search order.

```

(subclass Table Furniture)
(subclass DiningTable Table)
(subclass DiningTable Furniture)
(=>
  (instance ?P TransitiveRelation)
  (=>
    (and
      (?P ?A ?B)
      (?P ?B ?C))
      (?P ?A ?C)))

(instance subclass TransitiveRelation)

```

Figure 3.7: Example redundant knowledge base

This means that although two tree-structured proofs may be different, they may work out to be identical when redundant paths are removed and a linear proof structure generated. Sigma supports an option for hiding such duplicate proofs. One additional option is to suppress proofs that may have different steps but which lead to the same answer. Sometime it is useful to a knowledge engineer to see these alternate proofs, and sometime not.

Proofs are handled in the classes `ProofProcessor` and `ProofStep`.

3.7 Mapping, Merging and Translation

In addition to SUO-KIF and TPTP syntax, Sigma can also read and write OWL format [Bechhofer et al., 2004]. Since many lightweight ontologies are currently being created in OWL, this feature opens up the use of Sigma to a large community, and provides a straightforward migration path to use of a more expressive logic and more sophisticated inference. It also opens up the use of SUMO to a community that wishes to have simple and fast inference, since SUMO can be (and is) exported with a lossy translation to an OWL version. While the bulk of the SUMO axioms are not directly expressible in OWL, they can serve as informative comments (and in fact are exported as human-readable comments) that serve to better define terms for the human user than if they were simply omitted.

Sigma also includes an export of facts in Prolog form. Once Sigma generates a TPTP version of an ontology, the TPTPWorld tools also handle a translation to Prolog that supports horn clause rules. There is also a simple prototype capability for exporting SQL statements for database creation and population from Sigma.

The growing availability and coverage of lightweight taxonomies that cover

domain specific knowledge, and the corresponding phenomenon of "linked data" as a community objective has encouraged the addition of an ontology mapping and merging capability to Sigma. It is based on earlier work on a stand-alone tool [Li, 2004]. In mapping SUMO to simple taxonomies there is often very little information for the machine to use to determine what matches might exist. The principal problem appears to be massive numbers of false positive matches. A simple algorithm appears to do as well in practice as a more sophisticated one, since the bulk of effort is still spent by a human in selecting accurate matches. Having a simple and easy user interface appears to provide more leverage than an incrementally better matching algorithm.

3.8 Working with Sigma and SUMO

There are as many possible processes for formal ontology development as there are for software development. Small projects may benefit from the low overhead of an informal process. Large projects with big teams will benefit from a greater degree of formal process. A typical process employing Sigma to extend SUMO is as follows:

- Developers use a set of instructions or documents as a source, or write down text in natural language that describes the domain of interest.
- The text is used as a basis for creating a glossary of natural language terms and definitions. Developers examine the SUMO hierarchy (using the term browser, and tree/graph browser) for each term in the glossary. The WordNet search pages are used to find all the different meanings of each defined word in the source text, and the WordNet-SUMO mappings are used to find the formal SUMO term that best fits the intended meaning of the textual term. For any substantially new and specialized domain, the task is to find a more general term that encompasses the meaning of the more specific textual term. Textual terms that are already covered by specific SUMO definitions are put aside as complete. For new terms and definitions that are needed, developers begin by adding subclass or instance statements to the appropriate leaf term in SUMO, by creating and editing a text file in SUO-KIF format.
- Once a preliminary SUO-KIF file has been created, developers load it into Sigma, along with SUMO proper and all the other domain ontologies the new file may extend. Developers run the Sigma Diagnostics to find any errors.
- Developers use the information in the natural language definitions created earlier to guide creation of SUO-KIF axioms. Each class should have at least a subclass statement and a documentation statement. Each relation should have domain statements defining the class membership of its arguments, and be defined as an appropriate type of relation, such as

TransitiveRelation. Each term should have at least one rule, that helps to make the term usable for inference. If there are very few things that can be stated about the term, reconsider whether it should be created.

- Developers create **format** and **termFormat** statements in the language of choice to support natural language paraphrases in Sigma for the axioms previously written. These can be presented to domain experts to help confirm that the desired knowledge has been captured correctly. Developers map the terms in the ontology to WordNet. This is accomplished by placing links in the existing SUMO-WordNet mapping files (if mapping to English) that update the existing links where needed to point to the more specific terms that have just been created. Developers load the revised WordNet mapping files into Sigma and use the Sigma WordNet Diagnostics to see where the WordNet hierarchy may differ from the formal relationships created in the new ontology. The existence of differences is not necessarily bad, but they should be examined and understood.
- Developers run the Sigma Consistency Check to find any logical contradictions in the new theory. Normally, there will be many cycles of adding content, then running the Diagnostics and Consistency Check processes in Sigma to find and correct errors as the theory is elaborated. At each iteration where no errors are found, in a group development process, the theory would be uploaded to a source configuration management system such as CVS or Subversion. Other developers are then free to test new theories with respect to their own work, and coordinate with each other. One can view the Diagnostics and Consistency Check steps as analogous to compilation and build of a conventional procedural computer program.
- Peer review is one of the best ways to improve a theory. Sigma helps developers significantly beyond just reviewing a file of declarative code, allowing them to search and test a theory in many different ways.

3.9 Browsing and Editing

The Sigma system contains several different sorts of browsers for viewing formal knowledge bases. The most basic component is a term browser, which presents all the statements in which a particular term appears. The statements are sorted by argument position and then the appearance of the term in rules and non-rule statements are then shown. All the statements are hyperlinked to the terms that appear in them. The overall format of the browser page is handled in `Browse.jsp` (and `BrowseBody.jsp`, `BrowseHeader.jsp` and `BrowseExtra.jsp`). The back end code to produce the lists of axioms is in `HTMLformatter` and the most prominent method is `browserSectionFormatLimit()`. The list of words related to a particular SUMO term are produced in `WordNetUtilities.formatWordsList()` for English and if the user has selected another language from the Sigma language menu then `OMWordnet.formatWords()` will be called.

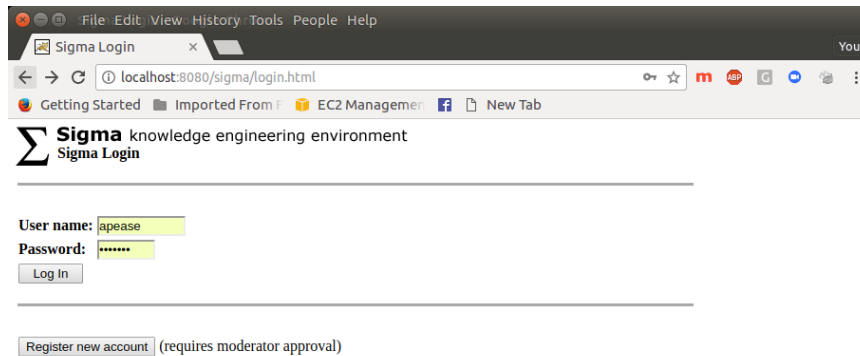


Figure 3.8: Sigma login screen

Two types of tree browser are provided. One provides an automatic graph layout and this is handled by the **Graph** class by producing a GraphViz¹ graph specification that it renders into an image file that gets displayed in the browser. Another shows a textual hierarchy. This is handled by the **TaxoModel** and **TaxoNode** classes. The user can choose the term to start with, the number of "levels" that should be presented from the term, and the binary relation to choose as the predicate that links the different nodes in the graph. For example, if the user asks for a graph of the term **IntentionalProcess** and for the subclass relation, the system will go "up" the graph to display the term **Process**, and down the graph to display the terms **Keeping**, **Guiding**, **Maintaining** etc. The user can ask for more levels up or down the graph. Note that any relation can be chosen, so, for example, a presentation of partonomies, or attribute hierarchies is also supported by choosing the relations **part** and **subAttribute** respectively.

3.10 Sigma Tutorial

In this section, we first run through a typical session illustrating major Sigma functions, and then provide a reference to the remaining functions. As discussed previously, KBs in first-order logic are similar to modern programming languages in complexity and expressiveness, and the most suitable tool for editing is a powerful text editor, such as one of the open source variants of Emacs, or a programming environment like Eclipse. The color-coding and formatting tools offered in such editors can be very helpful. Sigma serves the same purpose as a modern IDE, supporting structured examination, project management, and debugging.

¹<http://www.graphviz.org>

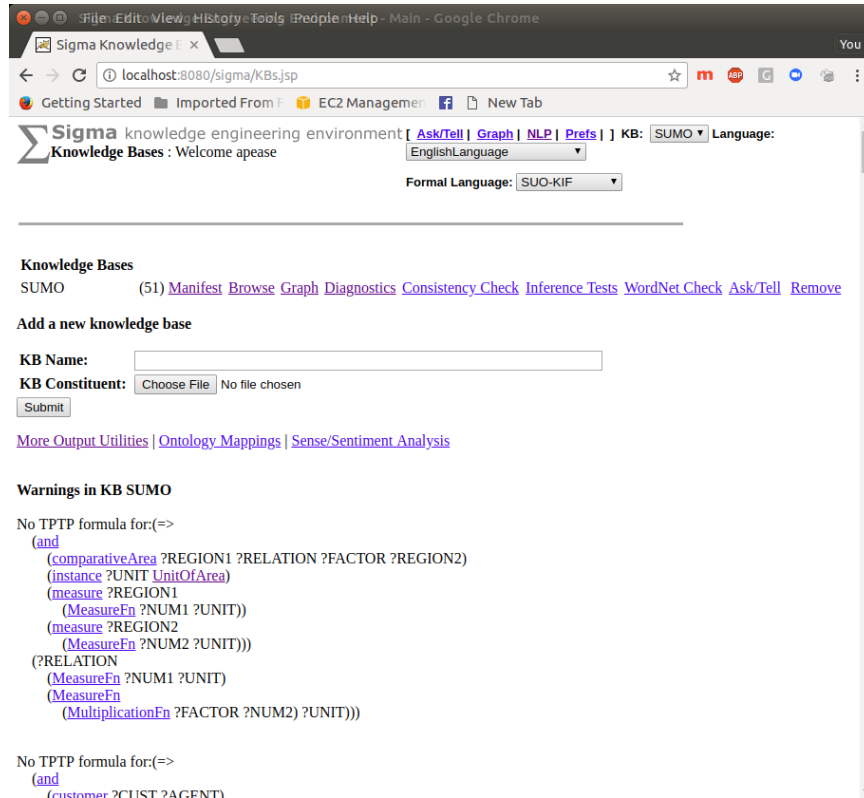


Figure 3.9: The Knowledge Bases screen

The first screen one sees is the login screen depicted in Figure 3.8. In its present version, Sigma has only the most rudimentary login functionality, with one hard-wired password that allows access to all Sigma functions. If that password is not chosen, only read-only operations are allowed. If you do not have the administrative password to Sigma and wish to perform operations that are not read-only, log in with the user name admin and the password admin.

The next screen, depicted in Figure 3.9, displays a listing of all the composite KBs loaded, the operations allowed on those KBs, and a type-in area for creating a new composite KB. The standard Sigma release package includes SUMO, so you should see a screen similar to that shown in Figure 3.9 if you have installed Sigma with the installation script, InstallSigma. If you are not using the standard release package, you will probably see a message indicating that no KBs are loaded. You will have to log in to Sigma with the administrator password and create a new knowledge base.

Sigma organizes ontologies in knowledge bases that are collections of files selected by the user. Each knowledge base consists of one or more constituents, which are files of statements written in SUO-KIF [Pease, 2009]. Each knowledge

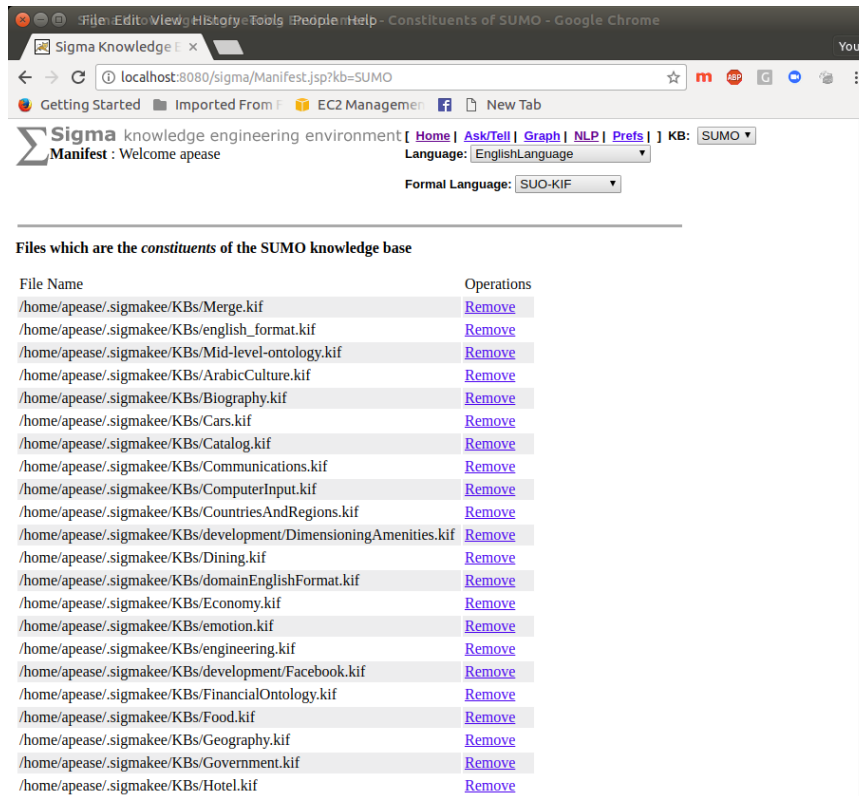


Figure 3.10: Manifest

base has a manifest which shows the files it contains (Figure 3.10).

At the bottom of the Manifest page (Figure 3.11) there are also controls that allow you to refetch knowledge bases from Git, reload the constituents, add a new constituent and save in different formats, including OWL, Prolog and TPTP.

The fundamental interface component of Sigma is a statement browser that displays the logical statements in which a given term appears. Clicking "Browse" on the Knowledge Bases screen displays a browser page for the selected knowledge base. The initial page just lists some statistics for the knowledge base, and provides a type-in area for entering terms from the knowledge base. Figure 3.12 shows a browser screen in which the user has typed the term **Walking**. Figure 3.13 shows a browser screen that lists all of the statements in the knowledge base in which **Walking** appears.

Clicking on a hyperlinked term in a statement displays the browser page for that statement. Clicking on the term **Running**, for example, causes the browser to show all statements in which **Running** appears. The browser also shows, in the blue-grey center column, the name of the KIF file from which the statement

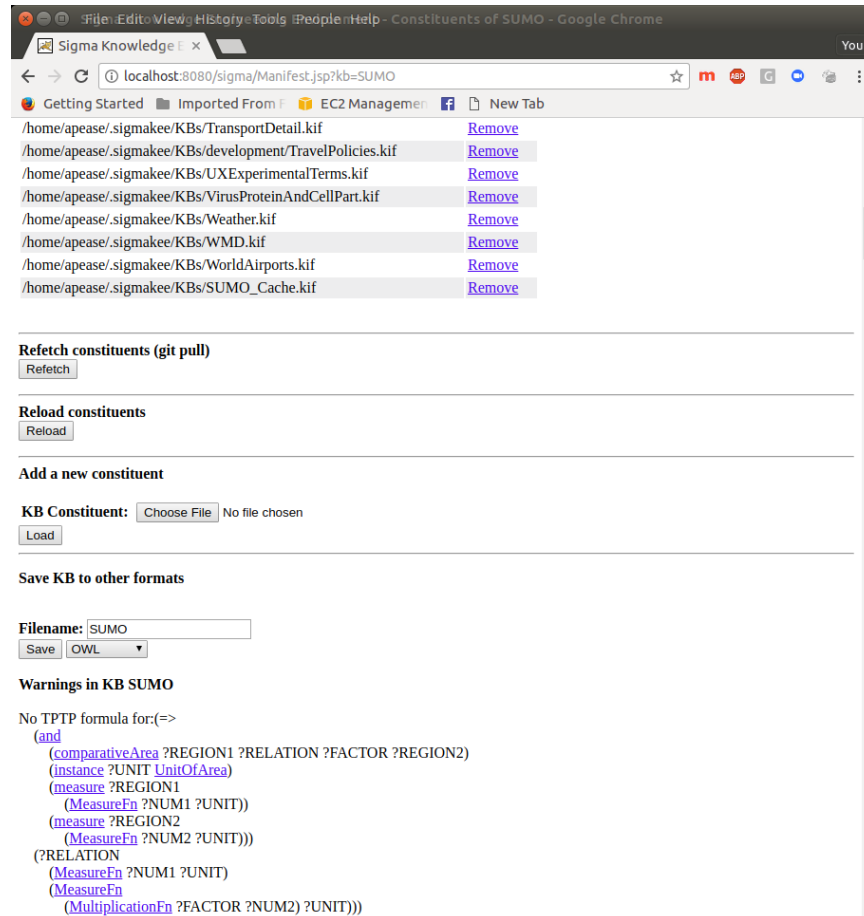


Figure 3.11: Manifest2

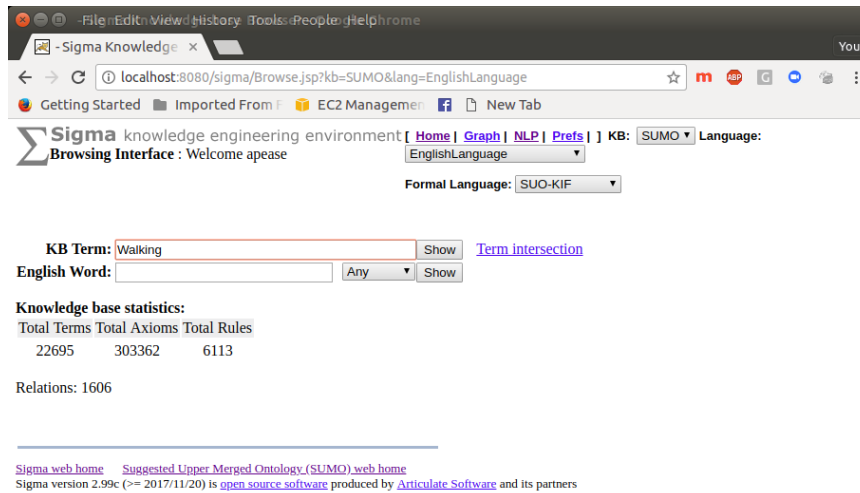


Figure 3.12: Initial term browser screen

was loaded, and the line number at which the statement appears in the file. An English paraphrase of the statement is shown in the right-hand column. The paraphrases are generated automatically from a set of format statements. On the Manifest page (Figure 3.10), we can see that the file `english_format.kif` has been loaded. That file contains language paraphrasing statements for English. In the term browser screens, one can see at the top right a pull-down menu labeled "Language". That menu is constructed automatically, based on the presence of language formatting statements in the knowledge base.

If the user types a word into the KB term box that cannot be found, Sigma responds with a set of terms which are closest, alphabetically, to the given term, as shown in Figure 3.14. By entering a word in the type-in area labeled "English Word", one can get a list of all the matching English word senses in the WordNet lexicon and their mappings to terms in SUMO [Niles and Pease, 2003]. The result of entering the word "buffalo", for example, is the page shown in Figure 3.15.

Another important functionality in Sigma is logical inference over KB contents. To explore this functionality, go to the Knowledge Bases screen and click on the link labeled "Ask/Tell". For a trivial inference, we can ask for a subclass of the class `Entity` by posing this query:

```
(subclass ?X Entity).
```

The inference engine returns a very simple proof that the class `Abstract` is a subclass of `Entity`, as shown in Figure 3.18.

Each step in the proof is numbered, and each step also has a justification for how it was derived. The justification can be a list of numbers. The value in the


Sigma knowledge engineering environment [[Home](#) | [Graph](#) | [NLP](#) | [Prefs](#)] | KB: [SUMO](#) Language: [EnglishLanguage](#)

Formal Language: [SUO-KIF](#)

KB Term: Show [Term intersection](#)

English Word: Any Show

Walking (walking) [\[assert facts\]](#)




[Rollerblade](#), [afoot](#), [amble](#), [ambulate](#), [ambulation](#), [angry walk](#), [break](#), [bumble](#), [canter](#), [career](#), [circumambulate](#), [clamber](#), [climb](#), [climb up](#), [clomp](#), [clump](#), [cock](#), [coogle](#), [constitutional](#), [constitutionalize](#), [countermarch](#), [crab](#), [creep](#), [curvet](#), [debouch](#), [dodder](#), [dostrot](#), [drag](#), [dressage](#), [drift](#), [err](#), [escalade](#), [exhibit](#), [falter](#), [fast break](#), [file](#), [file in](#), [file out](#), [fire walking](#), [flounce](#), [flounder](#), [foot](#), [footer](#), [footslog](#), [footstep](#), [forage](#), [gait](#), [gallop](#), [gimp](#), [go up](#)...

appearance as argument number 1

(documentation Walking ChineseLanguage "这是相对较慢的 Ambulating , 即是以至少总是有一只脚和地面 保持接触的方式移动。")	Merge.kif 10984-10985
(documentation Walking EnglishLanguage "Ambulating relatively slowly, i.e. moving in such a way that at least one foot is always in contact with the ground.")	Merge.kif 10982-10983
(externalImage Walking "http://upload.wikimedia.org/wikipedia/commons/6/6f/Walk-Cycle.gif")	pictureList.kif 3278-3278 "http://upload.wikimedia.org/wikipedia/commons/6/6f/Walk-Cycle.gif" is a URL , depicting walking
(externalImage Walking "http://upload.wikimedia.org/wikipedia/commons/d/d2/Marcheur_en_comp%C3%A9tion.jpg")	pictureList.kif 3279-3279 "http://upload.wikimedia.org/wikipedia/commons/d/d2/Marcheur_en_comp%C3%A9tion.jpg" is a URL , depicting walking
(subclass Walking Ambulating)	Merge.kif 10981-10981 Walking is a subclass of ambulating

appearance as argument number 2

(partition Ambulating Walking Running)	Merge.kif 10974-10974 Ambulating is exhaustively partitioned into walking and running
(subclass Hiking Walking)	Offers.kif 294-294 Hiking is a subclass of walking
(subclass Wading Walking)	Mid-level-ontology.kif 710-710 Wading is a subclass of walking

Figure 3.13: Term browser page for the SUMO term *Walking*

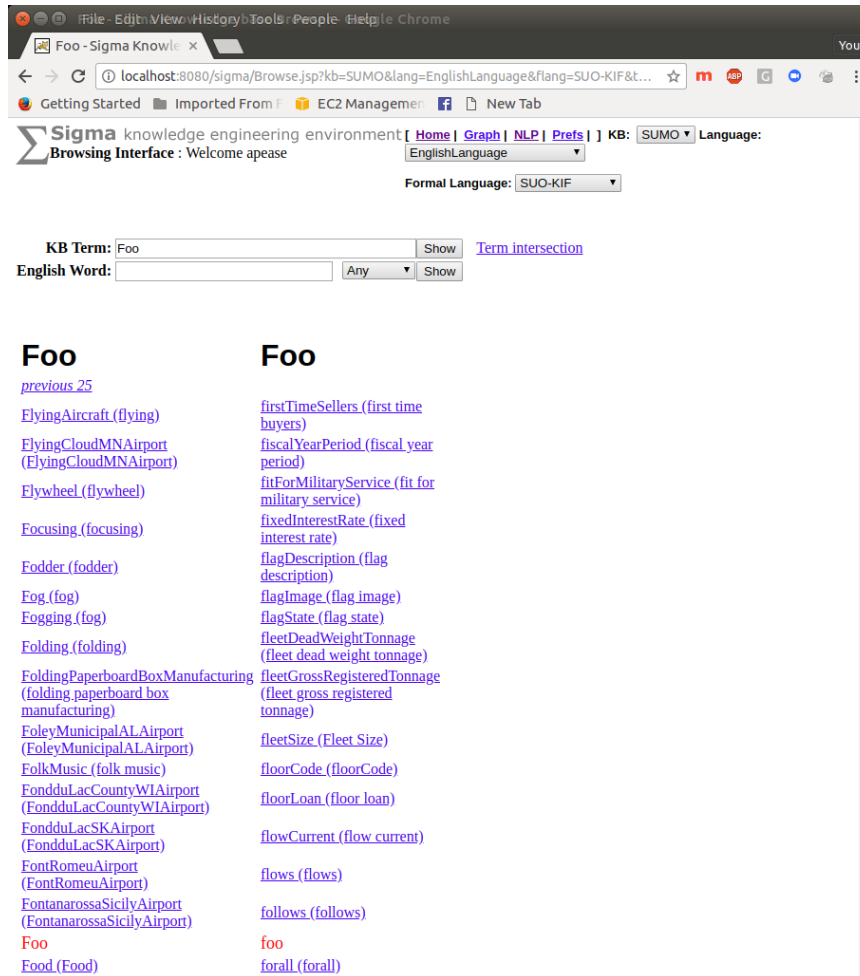


Figure 3.14: Term neighbors

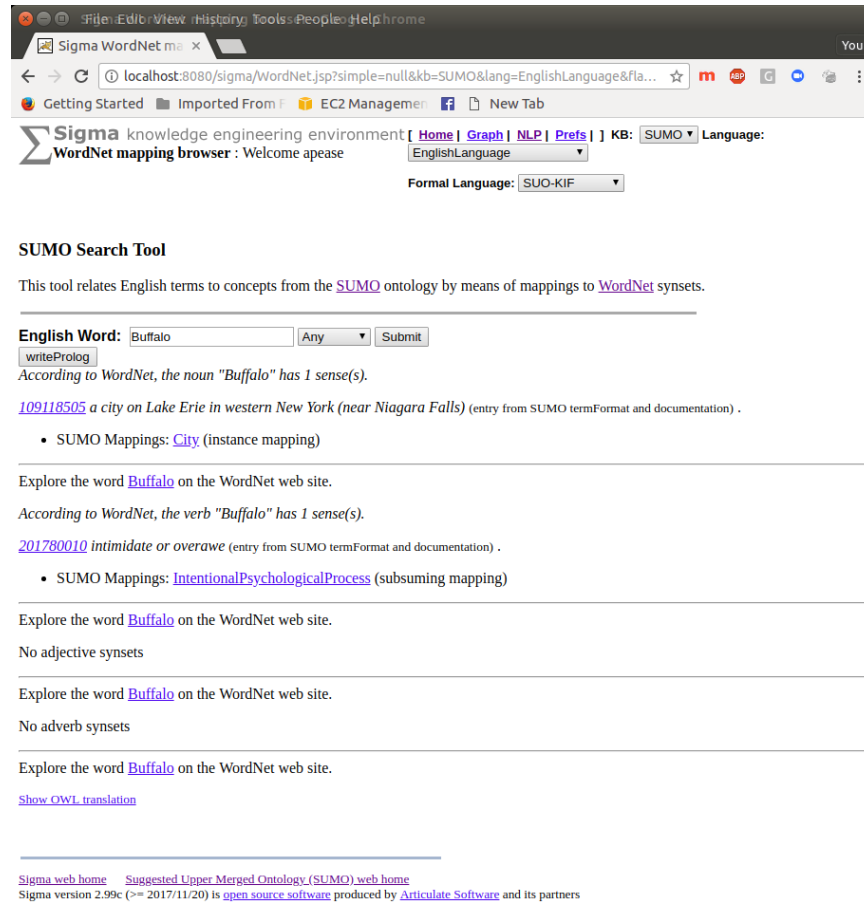


Figure 3.15: English word listing

Sigma KB Browser - Google Chrome Help

Sigma KB Browser x

localhost:8080/sigma/Graph.jsp?kb=SUMO&lang=EnglishLanguage&flang=SUO-KIF&rel... ☆ m ad C

Getting Started Imported From F EC2 Managemen f New Tab

Sigma knowledge engineering environment [Home | NLP | Prefs |] KB: SUMO Language: EnglishLanguage

KB Graph : Welcome apease

Formal Language: SUO-KIF

Relation: [subclass](#)

	direct-children	documentation	graph
Physical	16608	An entity that has a location in space-time. Note that locations are themselves understood to have ...	^
Object	13888	Corresponds roughly to the class of ordinary objects. Examples include normal physical objects, geo...	^
SphericalObject	.	An object with the attribute of being spherical.	^
Magnet	2	Any object that emits Magnetism . Note that this includes both NaturalMagnets and Electromagne...	^
TwoDimensionalObject	.	A real-world physical object with a very flat aspect. This includes drawings on paper, cave paintin...	^
TransitwayObstacle	3	TransitwayObstacle is the general class of Objects that can act as obstacles to Motion along ...	^
PreparedFood	80	PreparedFood refers to anything that undergoes some process intended to result in an object that ...	^
ContactSite	36	A ContactSite is an Object , typically a Place or a Residence or a CommunicationDevice suc...	^
SelfConnectedObject	10721	A SelfConnectedObject is any Object that does not consist of two or more disconnected parts.	^
Region	10551	A topographic location. Regions encompass surfaces of Objects , imaginary places, and Geograph...	^
Collection	657	Collections have members like Classes , but, unlike Classes , they have a position in space-tim...	^
Agent	1808	Something or someone that can act on its own and produce changes in the world.	^
AstronomicalBody	25	The Class of all astronomical objects of significant size. It includes SelfConnectedObjects lik...	^
Artifact	10736	An Object that is the product of a Making .	^

Relation: Term:

Levels "above": Levels "below": Total term limit: Show instances: ☒

Columns to display:

View format: ☐ graph ☒ text

[Sigma web home](#) [Suggested Upper Merged Ontology \(SUMO\) web home](#)

Sigma version 2.99c (>= 2017/11/20) is [open source software](#) produced by [Articulate Software](#) and its partners

Figure 3.16: Textual Graph page

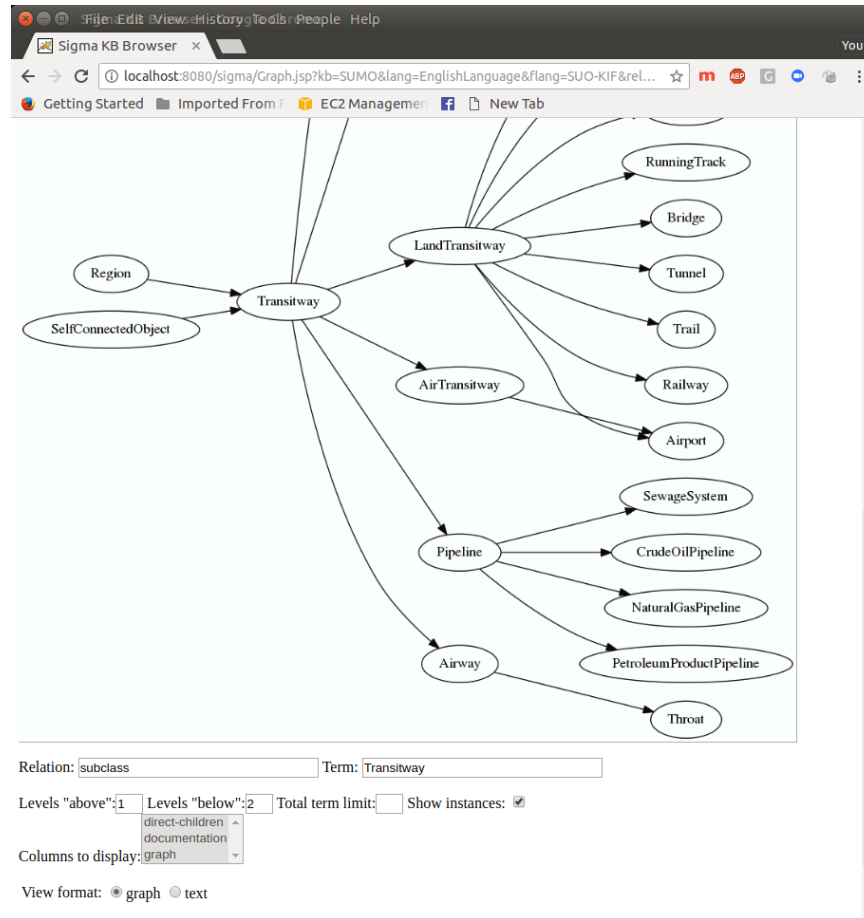


Figure 3.17: Graph page

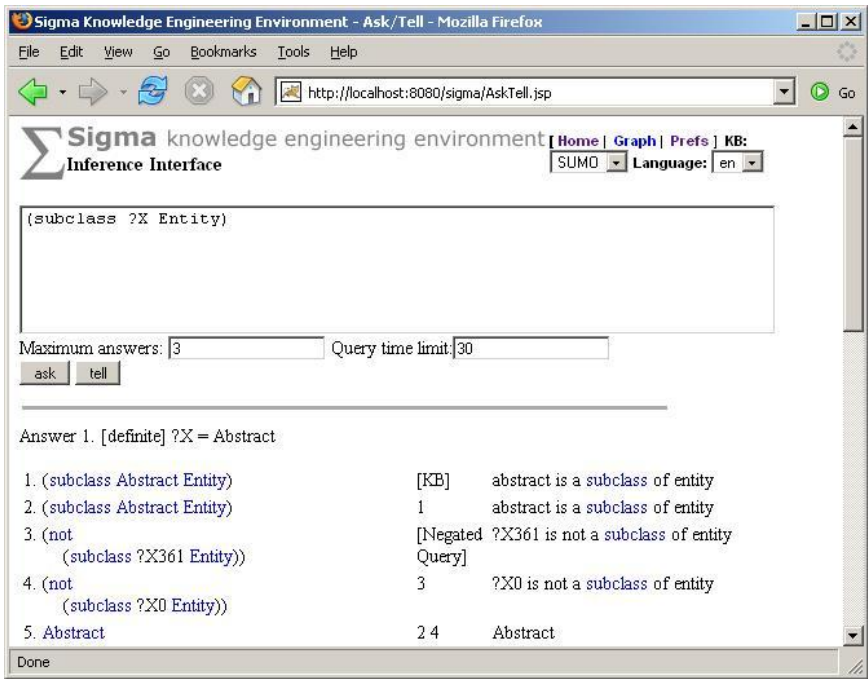


Figure 3.18: Sigma proof

justification column for step 4 shows that it was derived from step 3. Steps can also be taken directly from the knowledge base, which is denoted by "[KB]", or from the query itself. Currently, no further justification about the inference rule applied is provided, although in this case, one can see that step 4 is derived from step 3 simply by renaming the variable. The proof method employed is proof by contradiction, where the query is negated, and the system tries to find a contradiction that results. So, we see the label "[Negated Query]" as the justification for step 3. The inference engine can be controlled by limiting the number of answers it is directed to find, as well as by providing a time cutoff. On a large, interconnected knowledge base, there are so many possible search paths that the inference engine would frequently continue to search indefinitely if such cutoffs were not provided.

One can also assert an individual formula to the knowledge base by entering a KIF statement, and clicking the button labelled "tell". The asserted formula then becomes accessible for inference and browsing. If the user has not previously done a "tell" to the current KB, the system creates a new file called `<KB-name>_userAssertions.kif`, and adds the formula, as well as any subsequent formulas, to that file. The file will be loaded automatically when Sigma is restarted, and it can also be deleted from the manifest like any constituent, if desired.

Note that first-order inference is computationally expensive and expected results may not be achieved, even if they logically follow from the knowledge base. Also, the theorem prover has no notion of what it means to return a common-sense answer, just a logically correct one, so general axioms in SUMO can occasionally give rise to answers that, although logically true, may not be useful or expected.

Sigma has a simple function for caching subclass relationships that may improve many inferences, since reasoning about subclass relationships is often necessary. Sigma computes the transitive closure of subclasses statements and assert them directly, so that theorem proving does not have to apply the subclass reasoning axiom in SUMO during inference, but can find them as ground assertions. For example, if we ask whether `Human` is a subclass of `Object`, the theorem prover would have to apply the same axioms several times. It may, in fact, spend all of its allotted time exploring unhelpful search paths, since it does not know what common-sense answer we are looking for. By asserting directly that `Human` is a subclass of `Object`, we can short-circuit a number of spurious inference paths, and get better and faster results. Caching is turned on from the Preferences page (Figure 3.19), accessible via the "[Prefs]" link.

The Preferences page also allows the user to set the directory in which the inference engine executable is found, the address of the server running Sigma, the name of the KB containing SUMO, directory in which inference tests are found, and whether automatically cached statements should be displayed in the browser.

Sigma knowledge engineering environment [[Home](#) | [Graph](#) | [NLP](#) | [Prefs](#) |] KB: SUMO Language: EnglishLanguage
 Preferences : Welcome apease
 Formal Language: SUO-KIF

Fully qualified path and name of the inference engine
 Directory in which the CELT system is located
 Fully qualified path and name of SWI prolog
 DNS address of the computer on which Sigma is hosted
 Port number on which Tomcat responds
 Name of the SUMO KB in Sigma
 Directory in which tests for the inference engine are found
 Directory in which consistency check output can be found
 Directory in which local copy of TPTPWorld is installed
 Directory in which dot graphs will be saved
 Directory in which built in ATP systems are located
 Command to invoke text editor
 Command line option for text editor to set cursor at a particular line
 Directory where log files are saved.
 Level of messages to be logged (can be severe, warning, info, config, finest)

☐ yes ☒ no : Overwrite files of the same name when creating, copying, and loading KB constituents
☐ yes ☒ no : Show cached statements in the term browser
☐ yes ☒ no : Load CELT at startup

Formula Translation Options

1. ☒ yes ☐ no : Add a "sortal" antecedent to every axiom

Figure 3.19: Preferences page

Chapter 4

Conclusions

4.1 Acknowledgements

This manual would not be possible without the efforts of all the colleagues with whom I have written code and papers over the years.

- Professor Christoph Benzmler co-authored [Pease and Benzmler, 2010] and contributed the integration of LEO-II with Sigma and wrote the translation to THF.
- Michal Sevcenko developed the NLG language and a the KIFb SUMO browser that implemented it.
- Randy Schulz developed the first version of Sigma
- Geoff Sutcliffe and his students Steven Trac and Frank Theiss developed the TPTP interface.
- Stephan Schulz helped with the interface to Eprover
- Andrei Voronkov and Andrei Riazanov developed the interface to Vampire
- Andrei Holumb, Andrew Reed, Stelian Serban, Qingqing Cai, James Gung, Jerry Kurlandski, Piegen You, Harsh Kholi, Lyndon Xiang Nagaraj Bhat, Kyle Johnson, Gretchen Kern all contributed to Sigma during the year I spent at IPsoft.
- Thanks to Infosys (and Vishal Sikka, Alex Gray and Sanjay Mehta) for their generous support of this work, which allowed me to further develop the system and write this manual.

Bibliography

- [Bechhofer et al., 2004] Bechhofer, S., van Harmelen, F., Hendler, J., Horrocks, I., McGuinness, D., Patel-Schneider, P., Stein, L., Dean, M., and Schreiber, G. (2004). *OWL Web Ontology Language Reference*, *World Wide Web Consortium, Recommendation REC-owl-ref-20040210*. World Wide Web Consortium.
- [Deng et al., 2009] Deng, J., Dong, W., Socher, R., Li, L. J., Li, K., and Fei-Fei, L. (2009). (2009). ImageNet: A Large-Scale Hierarchical Image Database.
- [Fellbaum, 1998] Fellbaum, C. (1998). *WordNet: An Electronic Lexical Database*. Language, Speech, and Communication. MIT Press.
- [Hayes and Menzel, 2001] Hayes, P. and Menzel, C. (2001). (2001). A Semantics for Knowledge Interchange Format, in Working Notes of the IJCAI-2001 Workshop on the IEEE Standard Upper Ontology.
- [Hoder and Voronkov, 2011] Hoder, K. and Voronkov, A. (2011). Sine qua non for large theory reasoning. In *Proceedings of CADE-23*, Lecture Notes in Computer Science, pages 299–314, Wroclaw, Poland. Springer.
- [Hurd, 2003] Hurd, J. (2003). First-order proof tactics in higher-order logic theorem provers. In Archer, M., Vito, D., B., and Munoz, C., editors, *Proceedings of the 1st International Workshop on Design and Application of Strategies/Tactics in higher-order Logics*, number NASA/CP-2003-212448 in *NASA Technical Reports*, pages 56–68. NASA.
- [Li, 2004] Li, J. (2004). Lom: A lexicon-based ontology mapping tool. In *Proceedings of the 2004 Performance Metrics for Intelligent Systems conference (PerMIS)*, pages 321–325. NIST.
- [Niles and Pease, 2003] Niles, I. and Pease, A. (2003). Linking Lexicons and Ontologies: Mapping WordNet to the Suggested Upper Merged Ontology. In *Proceedings of the IEEE International Conference on Information and Knowledge Engineering*, pages 412–416.
- [Pease, 2003] Pease, A. (2003). The Sigma Ontology Development Environment. *Working Notes of the IJCAI-2003 Workshop on Ontology and Distributed Systems*, 71.

- [Pease, 2009] Pease, A. (2009). Standard upper ontology knowledge interchange format. http://sigmakee.cvs.sourceforge.net/*checkout*/sigmakee/sigma/suo-kif.pdf.
- [Pease, 2013] Pease, A. (2013). Sigma web site. <http://sigmakee.sourceforge.net>.
- [Pease and Benzmler, 2010] Pease, A. and Benzmler, C. (2010). Ontology Archaeology: Mining a Decade of Effort on the Suggested Upper Merged Ontology. *The ECAI-10 Workshop on Automated Reasoning about Context and Ontology Evolution*.
- [Pease and Li, 2010] Pease, A. and Li, J. (2010). Controlled English to Logic Translation. In Poli, R., Healy, M., and Kameas, A., editors, *Theory and Applications of Ontology*. Springer.
- [Pease and Sutcliffe, 2007] Pease, A. and Sutcliffe, G. (2007). First Order Reasoning on a Large Ontology. In *Proceedings of the CADE-21 workshop on Empirically Successful Automated Reasoning on Large Theories (ESARLT)*.
- [Pease et al., 2010] Pease, A., Sutcliffe, G., Siegel, N., and Trac, S. (2010). Large Theory Reasoning with SUMO at CASC. *AI Communications, Special issue on Practical Aspects of Automated Reasoning*, 23(2-3):137–144.
- [Riazanov and Voronkov, 2002] Riazanov, A. and Voronkov, A. (2002). The Design and Implementation of VAMPIRE. *Journal of AI Communications*, 15(2/3):91–110.
- [Russel and Norvig, 2009] Russel, S. J. and Norvig, P. (2009). *Artificial Intelligence: A Modern Approach*. Prentice Hall, 3rd edition.
- [Schulz, 2002] Schulz, S. (2002). E – A Brainiac Theorem Prover. *Journal of AI Communications*, 15(2/3):111–126.
- [Sevcenko, 2003] Sevcenko, M. (2003). Knowledge support for modeling and simulation. In *Proceedings of the 7th International Conference, Knowledge-Based Intelligent Information and Engineering Systems, KES 2003*, pages 99–103, Oxford, UK. Springer Berlin Heidelberg.
- [Sutcliffe, 2009] Sutcliffe, G. (2009). The TPTP Problem Library and Associated Infrastructure: The FOF and CNF Parts, v3.5.0. *Journal of Automated Reasoning*, 43(4):337–362.
- [Trac et al., 2007] Trac, S., Puzis, Y., and Sutcliffe, G. (2007). An interactive derivation viewer. *Electronic Notes in Theoretical Computer Science (ENTCS)*, 174(2).
- [Trac et al., 2008] Trac, S., Sutcliffe, G., and Pease, A. (2008). Integration of the TPTPWorld into SigmaKEE. In *Proceedings of IJCAR '08 Workshop on Practical Aspects of Automated Reasoning (PAAR-2008)*. CEUR Workshop Proceedings.

Index of Classes

CCheck, 43
CCheckManager, 43
Clausifier, 26

Diagnostics, 39

EProver, 40

Formula, 26
FormulaPreprocessor, 16, 20, 22, 26

Graph, 47

HTMLformatter, 46

InferenceEngine, 40

KBcache, 17
KBmanager, 13, 14
KIF, 14, 15

LanguageFormatter, 34

PredVarInst, 22
ProofProcessor, 44
ProofStep, 44

RowVars, 24

SInE, 40
SUMOformulaToTPTPformula, 26
SUMOKBtoTPTPKB, 26
SystemOnTPTP, 41

TaxoModel, 47
TaxoNode, 47
THF, 24

Vampire, 41

WNdiagnostics, 40
WordNetUtilities, 46
WSD, 18

Index

- arity, 22
- Browse, 49
- buffalo, 51
- Caching, 58
- CASC, 43
- Consistency Check, 46
- contradiction, 39
- CVS, 46
- debugging, 47
- description logic, 38
- diagnostics, 40, 45
- Eclipse, 47
- efficiency, 25
- Emacs, 47
- format, 46
- Git, 8
- Guiding, 47
- higher-order logic, 24
- IDE, 47
- ImageNet, 38
- InstallSigma, 48
- IntentionalProcess, 47
- Keeping, 47
- linked data, 45
- located, 23
- Maintaining, 47
- Manifest, 51
- Metis, 42
- part, 47
- partlyLocated, 23
- partonomy, 47
- peer review, 46
- predicate variable, 22
- Preferences, 58
- printf, 34
- Prolog, 28, 44
- proof, 25
- row variables, 23
- SInE, 40
- subAttribute, 47
- subrelation, 23
- Subversion, 46
- SUO-KIF, 7
- termFormat, 46
- TPTP, 7
- TPTPWorld, 20
- transitive binary relation, 38
- TransitiveRelation, 22, 46
- variable binding, 25
- Walking, 49
- Wikipedia, 38
- WordNet, 51